



100 Snowflake Cost Optimization Techniques

ALL ABOUT

COST AND PERFORMANCE OPTIMIZATION

IN SNOWFLAKE

100 Snowflake Cost Optimization Techniques

ALL ABOUT

COST AND PERFORMANCE OPTIMIZATION

IN SNOWFLAKE

Best Ways to Benefit from this Course

- * Free Trial Snowflake Account
- * VSCode Project Setup + GitHub
- * 10 Sections w/ Introductions
- * Eventual Hands-On + End-Review (slides)
- * Quizzes to Test Your Knowledge
- * Captions + Playback Speed
- * Q&A + Money Back Guarantee
- * Review

Introduction to Virtual Warehouses

Tip #1: Larger Virtual Warehouses May Actually Cost You Less

Tip #2: Auto-Suspend Any Warehouse After One Minute

Tip #3: Any Resumed Warehouse Will Cost You at Least One Minute

Tip #4: Never Auto-Suspend Any Warehouse After Less Than One Minute

Tip #5: X-Small Warehouses Could Be Powerful Enough

Tip #6: Resized Warehouses are for More Complex Queries

Tip #7: Multi-Cluster Warehouses are for Multiple Users and Concurrency

Tip #8: Multi-Cluster Warehouses Should Always Have Min Clusters 1

Tip #9: Use Economy Scaling Policy To Save Money


Tip #10: When to Use Snowpark-Optimized Warehouses

*Larger Warehouses
May Actually
Cost You Less*

Tip #1

WAREHOUSE SIZE VS QUERY EXECUTION TIME

■ Nodes ■ Seconds



Cost

4X-Large: 128 Nodes x 60 secs = 128 x 5c = **\$6.40**

5X-Large: 256 Nodes x 30 secs = 256 x 2.5c = **\$6.40**

6X-Large: 512 Nodes x 15 secs = 512 x 1.25c = **\$6.40**

at \$3/credit = \$3/hour/node = 5c/minute/node

Checkpoints

- * virtual warehouses bring 80-90% of all Snowflake costs!
- * **cost effective** = ~same cost, but using more resources
- * double in size WH + query twice faster → ~same cost!
- * bigger warehouses can actually save you money!
- * at some point, bigger warehouses will no longer execute so fast
- * small warehouses could take forever in very complex queries

*Auto-Suspend
Any Warehouse
After One Minute*

Tip #2

Cost

query execution time: 4 min/hour = 96 min/day x 5c = **\$4.80**

1 min auto-suspend: 4 min/hour = 96 min/day x 5c = **\$4.80**

5 min auto-suspend: 20 min/hour = 480 min/day x 5c = **\$24**

for one query taking one min every 15 min

w/ X-Small warehouse → \$3/credit = 5c/min (running!)

Checkpoints

- * idle warehouses cost money (the same!)
- * reducing auto-suspend to the min will save you big!
- * creating a warehouse will auto-resume it
- * can set `INITIALLY_SUSPENDED` to True only by code
- * avoid auto-suspend by executing more queries together

*Resumed Warehouses
Will Cost You
at Least One Minute*

Tip #3

Cost

suspend after 10 secs: 192 nodes x 1 min x 5c = **\$9.60**

suspend-resume in 10 secs: 192 nodes x 2 min x 5c = **\$19.20**

for 3 clusters w/ 3X-Large WH each = 3 x 64 nodes = 192 nodes

Checkpoints

- * you are charged at least one minute for any WH resume
- * min auto-suspend (recommended) is one minute
- * default auto-suspend is 5 minutes → change it to 1 min
- * a background task checks every minute for auto-suspend
- * the one-minute period restarts when warehouse resumed
- * suspend+resume within a minute may charge you twice

*Never Auto-Suspend
Any Warehouse
After Less Than a Minute*

Tip #4

Edit Warehouse

📁 WH_LARGE as 👤 ACCOUNTADMIN

+ Add Comment

Type: Standard ▾

Size: 3X-Large ▾ ?

Advanced Options ^



Auto resume

Automatically resumes the warehouse when any statement that requires a warehouse is submitted.



Auto suspend

Automatically suspends the warehouse if it is inactive for the specified period of time.

Suspend After

1

min(s) of inactivity

Cancel

Save Warehouse


Checkpoints

- * in Snowsight → min auto-suspend is one minute (recommended)
- * in SQL → min auto-suspend is 0 seconds (60 recommended)
- * auto-suspend zero means you always keep it alive
- * a background task checks every minute for auto-suspend
- * the one-minute period restarts when warehouse resumed
- * suspend+resume within a minute may charge you twice


*X-Small Warehouses
Could Be
Powerful Enough*

Tip #5

X-Small Warehouse (1 node)



Small Warehouse (2 nodes)



Checkpoints


- * a "node" is way more than a one-CPU VM
- * a single "node" is rather a powerful server w/ 8-vCPUs
- * do not overspend, as X-Small could be enough
- * assumed 16 GB RAM + 200 GB SSD
- * for AWS: c5d.2xlarge EC2 instance (except for 5XL and 6XL)
- * but only Snowflake knows what's in a node (CPU, RAM, SSD)

*Resized Warehouses
are for More
Complex Queries*


Tip #6

Fix Warehouse Size


Large
Warehouse



Large
Warehouse



Manual Scale





Checkpoints

- * bigger sizes could help more complex queries execute faster
- * bigger warehouse size could cost you less (as seen before)
- * no concurrency improvement by resizing
- * resizing a warehouse is always a manual process
- * one query can be executed by only one warehouse
- * cannot have different sizes in a multi-cluster warehouse


*Multi-Cluster Warehouses
are for Multiple Users
and Concurrency*

Tip #7

Fix Cluster Size



Auto Scale



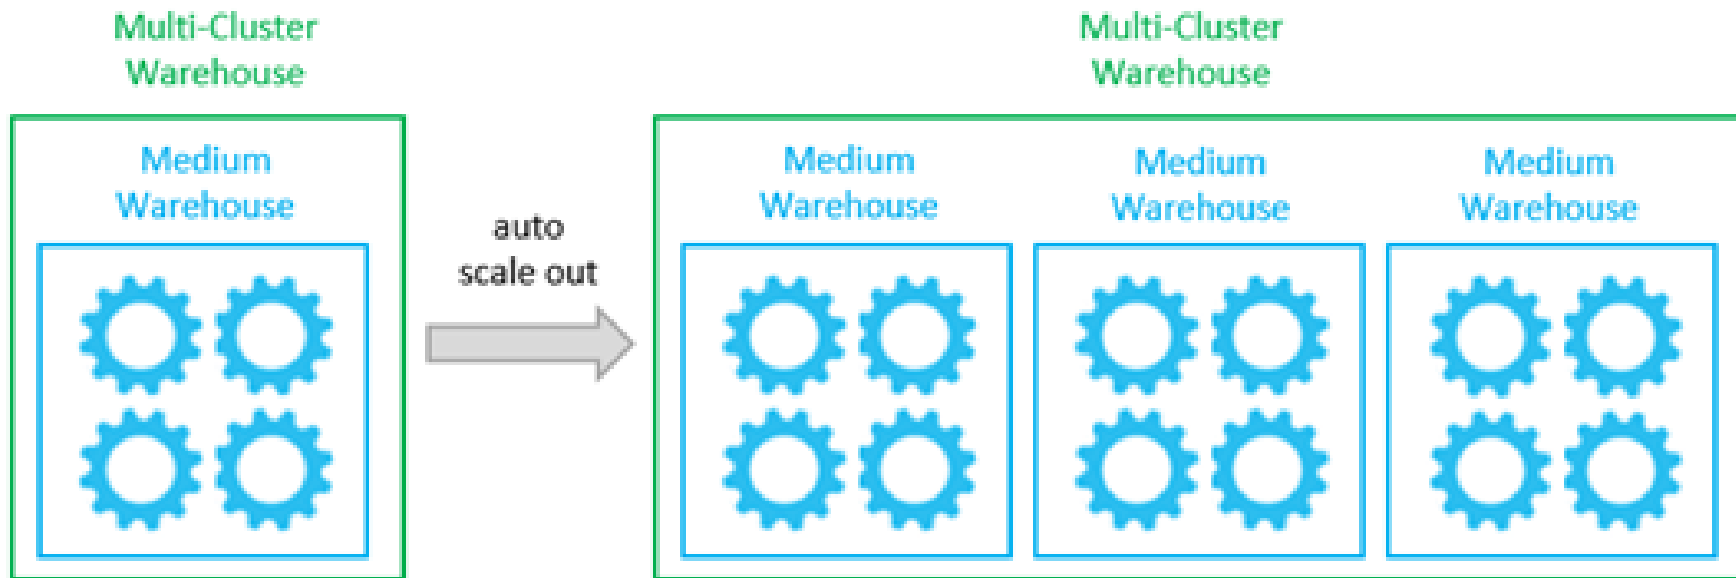
Checkpoints

- * multi-cluster warehouses improve concurrent access
- * multi-cluster warehouses do not help with complex queries
- * multi-clustering = main reason you should have Enterprise Edition
- * cannot have different sizes in a multi-cluster warehouse
- * could be avoided by assigning diff users/queries to diff warehouses
- * support auto-scaling (between 1..10 clusters)
- * one query can be executed by only one warehouse

*Multi-Cluster Warehouses
Should Always Have
Min Clusters 1*

Tip #8

Auto Scale



Checkpoints

- * do not over-provision w/ min cluster count > 1
- * unnecessary costs
- * you may have one user running queries w/ very low latency (no queuing)
- * min cluster count = max cluster count \rightarrow no auto-scale

*Use Economy
Scaling Policy
To Save Money*

Tip #9

Multi-Cluster Warehouses

* *scale modes*

- **maximized** → min=max clusters
- **auto-scale** → between min < max clusters

* *scaling policies*

- **standard** (def) → for end-user live queries (performance is top priority)
- **economy** → for lower priority batch jobs (to optimize cost)

Standard Scaling Policy

- * *auto-resume*: after 20 secs when new queued req + prior WH at its capacity
- * *auto-suspend*: after 2-3 checks every min, if queries on least loaded WH can be distributed elsewhere
- * avoid query queuing without saving credits
- * for end-user live queries, when performance is a priority

Economy Scaling Policy

- * *auto-resume*: when crt queue load could keep new WH busy for min 6 mins
- * *auto-suspend*: after 5-6 checks every min, if queries on least loaded WH can be distributed elsewhere
- * save credits, query queuing may happen
- * for lower priority batch jobs, to maximize throughput + optimize cost

Checkpoints

- * **standard scaling policy** (def) = will avoid queuing, but save no credits
- * **economy scaling policy** = will save credits, but have query queuing
- * auto-scale in (WH shutdown) could take 5-6 or 2-3 minutes!
- * avoid scale mode *maximized* → prefer *auto-scale* between min..max
- * avoid **MAX_CLUSTER_COUNT** 10 → keep 5 ok
- * **MAX_CONCURRENCY_LEVEL** (def 8) = queries in parallel on a WH

*When to Use
Snowpark-Optimized
Warehouses*

Tip #10

Costs

Type of Virtual Warehouse	X-Small	Small	Medium	Large	X-Large	2X-Large	3X-Large	4X-Large	5X-Large	6X-Large
Standard	1	2	4	8	16	32	64	128	256	512
Snowpark-optimized	n/a	n/a	6	12	24	48	96	192	384	768

Checkpoints

- * they have 16x more memory than Standard WHs + cost 33% more
- * they are more for Data Science and ML experiments
- * for one stored proc deployed by Snowpark in a single-node WH
- * well-suited for model training in ML (not model serving!)
- * no GPUs → see rather Snowpark Container Services w/ NVIDIA

Introduction to Compute Workloads

Tip #11: Use Resource Monitors

Tip #12: Use Account-Level Budgets

Tip #13: Prevent Never-Ending Queries

Tip #14: Manually Kill Running Queries

Tip #15: Reduce Warehouse Sizes

Tip #16: Consolidate All Warehouses

Tip #17: Use Parallel Jobs for Batch Transformations

Tip #18: Avoid Checking Too Much on Metadata

Tip #19: Charts for Warehouse Monitoring

Tip #20: Revisit the Main Traps with Warehouses

Use

Resource Monitors

Tip #11

CREATE RESOURCE MONITOR

- * *levels*: account / warehouse ← SET RESOURCE_MONITOR = ...
- * with *credit_quota* ← auto-reset monthly by default
- * *triggers* (on N percent do) notify / suspend / suspend_immediate
- * *frequency* = monthly / weekly / daily / yearly / never
- * *start_timestamp* (= immediately) ... *end_timestamp*
- * *notify_users* = ("CRISTISCU", ...) ← all ACCOUNTADMIN by def

Checkpoints

- * specify one monitor for the whole account
- * specify one monitor for each virtual warehouse (1-N)
- * enable notifications for each account admin
- * watch all warehouses with no res mon assigned
- * suspend immediate if long running queries expected
- * schedule, for batch queries w/ periodic runs

*Use
Account-Level
Budgets*

Tip #12

Checkpoints

* *account budget*

- PuPr, only for paying accounts, must enable/activate first
- max monthly spending limit → auto-send emails when exceeding
- SNOWFLAKE.LOCAL.ACCOUNT_ROOT_BUDGET class

* *custom budgets*

- on group of resources, in db.schema, max 100
- dbs, schemas, tables, mat views, WHs, pipes, tasks, serverless
- SNOWFLAKE.CORE.BUDGET class

*Prevent
Never-Ending
Queries*

Tip #13

Checkpoints

- * a running query is canceled by the system after 2 days
- * **STATEMENT_TIMEOUT_IN_SECONDS** = sess/WH/acct
 - def 2 days (172,800 secs), 0 to 604800 (7 days)
 - max overall time = queue + locked + exec + compile + ...
 - also for queries w/o WH (DDL/DML: CLONE, CREATE...)
- * **STATEMENT_QUEUED_TIMEOUT_IN_SECONDS** = def 0

*Manually Kill
Running Queries*

Tip #14

Checkpoints

- * *Cancel* button in Snowsight, X button in Query History
- * **SYSTEM\$CANCEL_QUERY** with last_query_id()
- * CTRL+C from SnowSQL
- * only own queries (not necessarily ACCOUNTADMIN)
- * NOT with: close browser, abort session, suspend WH
- * could also manually stop WH when query ends, to save money!

*Reduce
Warehouse Sizes*

Tip #15

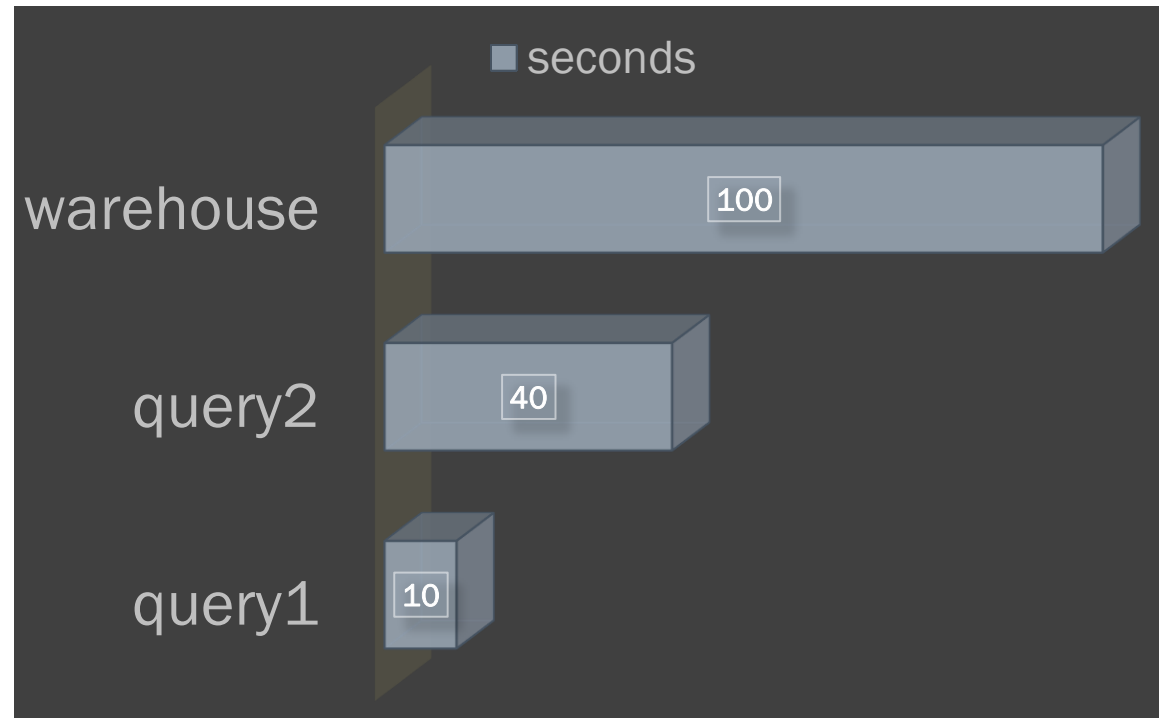

Checkpoints

- * start with the smallest 1-node X-Small warehouse
- * start with MIN_CLUSTER_COUNT 1 in multi-cluster WH
- * as a user, connect with an X-Small default WH
- * upsize only when need for larger WH is proven
- * use Snowpark-optimized only for mem-intensive workloads
- * drop expensive warehouses, not frequently used
- * cancel queries taking too long w/ current warehouse

*Consolidate
All Warehouses*

Tip #16

Query Execution



Warehouse Load/Utilization

- * *query1 load*: $10/100/2 = 5\%$ (could run 20 such queries in 100 secs)
- * *query2 load*: $40/100/2 = 20\%$ (could run 5 such queries in 100 secs)
- * *warehouse load/utilization*: $0.05 + 0.2 = 25\%$

query1 (10 secs) + query2 (40 secs)

warehouse `max_concurrency_level` = 2

warehouse up: 40 secs (queries) + 60 secs (auto-suspend) = 100 secs

System Views

- * **WAREHOUSE_METERING_HISTORY**

- *start_time .. end_time + warehouse_name*
- *credits_used*

- * **QUERY_HISTORY**

- *start_time .. end_time + warehouse_name*
- *execution_time + query_load_percent*


Checkpoints

- * **consolidation** = reduce # of *underutilized* warehouses (cost money!)
- * *overutilized* warehouses (w/ queuing) have no cost impact
- * *query latency* (queuing) vs *warehouse utilization* ← compromise!
- * *query load* vs *warehouse load* (+ MAX_CONCURRENCY_LEVEL)
- * redirect queries from some WHs to other WHs ← see utilization heatmap
- * group similar workloads in the same WH
- * *workload type (util. target)*: ETL (95%), interactive (80%), dashboards (60%)

*Use Parallel Jobs
for Batch Transformations*

Tip #17

Data Transformation Pipeline



Checkpoints

- * ~parallel ingestion into table from staged files
- * split intensive MERGE statements w/ filters
- * send multiple SQL jobs in parallel on diff connections
- * max concurrency on 1-node X-Small WH is 10
- * increase WH size for better distribution
- * multi-cluster WH could also make sense for >> ops

*Avoid Checking
Too Much
on Metadata*

Tip #18

Snowsight

Cost Management

• COMPUTE_WH

Account Overview

PREVIEW

Consumption

Resource Monitors

Account spend for **KHB53464** ⓘ from **Feb 26 - Mar 4** ▾

\$ 5.27 —
15.68 Spend in
Spend in credits
USD

Compute
price/credit

\$1.96 0.66
ⓘ Average daily Average daily
cost credits

Checkpoints

* *Snowsight*

- Snowflake's admin console will incur a cost!
- check Snowsight screens w/ a WH
- auto-refresh Query History every N seconds

* *querying metadata*

- querying Information Schema does not require a WH
- querying Account Usage schema requires a WH
- many Account Usage queries are rather slow (ex: *grants_to_roles*)
- many Account Usage views could return a lot of data (ex: *query_history*)
- many columns return JSON data, that needs flattening (ex: *access_history*)

Charts for Warehouse Monitoring

Tip #19

Checkpoints

- * overall compute: warehouses / cloud services / serverless
- * WAREHOUSE_METERING_HISTORY view
- * top warehouses by cost → credits used
- * utilization heatmaps → warehouse loads
- * queries → running/provisioning/blocked/queued
- * total spending per user / query type ...
- * timeline charts

*Revisit
the Main Traps
with Warehouses*

Tip #20

Checkpoints

- * no predef. resource monitors + not through UI → no proactive approach to overspending
- * `AUTO_SUSPEND` by def 600 (10 minutes) in code + 5 min in UI
- * billed per second, but at least 1 min + `AUTO_SUSPEND` 1 min at least (can be set lower!)
- * manual 2+ resume+suspend pairs in the first minute → billed for 2+ minutes!
- * `INITIALLY_SUSPENDED` by def FALSE and no UI access → bill 1 min no matter what
- * `MIN_CLUSTER_COUNT` could be set to >1 → should be always 1
- * `STATEMENT_TIMEOUT_IN_SECONDS` 2 days by def + hidden Cancel btn in Query History
- * `STATEMENT_QUEUED_TIMEOUT_IN_SECONDS` 0 by def (no timeout)

Introduction to Snowflake Accounts

Tip #21: What to Choose for a Free Trial Account

Tip #22: When to Use a Free Trial Account

Tip #23: Understand Price Tables for Warehouse Compute Services

Tip #24: Understand Price Tables for Cloud and Serverless Services

Tip #25: Understand Price Tables for Storage and Data Transfer

Tip #26: Use the Account Overview Interface in Snowsight

Tip #27: Use Organization Accounts

Tip #28: Limit Warehouse Changes with Access Control

Tip #29: Adjust Default Values of Account-Level Parameters

Tip #30: Careful with Reader Accounts

*What to Choose
for a Free Trial Account*

Tip #21

signup.snowflake.com

Cristian

✓

Scutaru

✓

your-email@gmail.com

✓

XtractPro Software

✓

Data Engineer

▼

Canada

▼

☒ No, I do NOT want Snowflake to send me e-mails about products, services, and events that it thinks may interest me.


Choose your Snowflake edition*


☐ **Standard**
A strong balance between features, level of support, and cost.

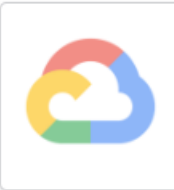
☒ **Enterprise**
Standard plus 90-day time travel, multi-cluster warehouses, and materialized views.

☐ **Business Critical**
Enterprise plus enhanced security, data protection, and database failover/fallback.

Choose your cloud provider*


Microsoft Azure


Amazon Web Services


Google Cloud Platform

US West (Oregon)

▼

Thanks for signing up with Snowflake!

Help us better serve you by answering these questions.

Why are you signing up for a trial?

☐ Company is considering Snowflake

☐ Training or Certification

☐ Personal learning and development

☐ Other

Skip

Continue

Checkpoints

- * no credit card required / no obligations whatsoever
- * email address could be reused after a while
- * no need for a corporate email address
- * default Enterprise edition could be the most appealing
- * AWS is frequently used first as provider for new PuPr features
- * region closest to where you are / where your data is
- * can skip the last survey questions

*When to Use
a Free Trial Account*

Tip #22

Limitations

- * running out of email addresses
- * some new features (Budgets) not on free trial accounts
- * native apps sharing requires accounts within org
- * no Snowflake customer service
- * you'll need to upload/reinstall all after one month
- * you lose all access after one month

Checkpoints

- * trial of native apps from Snowflake Marketplace
- * experiments on large public datasets
- * prototyping + trying new Snowflake features
- * introducing a new potential client to Snowflake
- * prepare for + write a new Udemy course

*Understand Price Tables
for Virtual Warehouse
Compute Services*

Tip #23

VW Compute Pricing

- * **Standard Edition** = \$2/credit/node = **\$2**/h for X-Small WH
- * **Enterprise Edition** = \$3/credit/node = **\$3**/h for X-Small WH
- * **Business Critical Edition** = \$4/credit/node = **\$4**/h for X-Small WH

in AWS/US region, for active 1-Node WH

Small (2-Nodes) is 2x more expensive, Medium (4-Nodes) 4x etc.

Snowflake Pricing

- * **Compute** = w/ credits, for active **VW/SPCS**, per second, minimum 1 min
- * **Cloud Services** = 4.4 credits/h, not billed if < 10% of daily VW credits
- * **Serverless Features** = on managed compute (+cloud services), per second
- * **Data Sharing** = pay for all Reader Accounts, not for other Snowflake accounts
- * **Storage** = avg TB/mo (snapshots of data), compressed, diff for hybrid tables!
- * **Data Transfer** = Ingress/Egress, separate for external functions

Checkpoints

- * *all compute*

- VW Compute
- Cloud Services
- Serverless Features

- * *for the same supported queries*

- Enterprise is 1.5x more expensive than Standard
- Business Critical is 2x more expensive than Standard

*Understand Price Tables
for Cloud and
Serverless Services*

Tip #24

Cloud Services

- * distributed back-end systems
- * SHOW commands
- * query compilation
- * result caches
- * INFORMATION_SCHEMA queries

Serverless Features

- * [Query Acceleration Service \(QAS\)](#) - Nx boost of a WH
- * [Search Optimization Service \(SOS\)](#) - (also in secondary db)
- * [Automatic Clustering](#) - for clustered tables (w/ clustered key), initial + recluster
- * [Materialized Views](#) - for maintenance (also in secondary db), auto-refresh data
- * [Snowpipe & Snowpipe Streaming](#) - w/ fixed credits/files or h
- * [Serverless Tasks](#) - 1.5x Managed Tasks (w/ WH)
- * [Hybrid Tables](#) - for maintenance+requests, indexes, R/W data
- * [Replication](#) - only in Business Critical for objects other than dbs

Checkpoints

* *cloud services*

- charged only when $> 10\%$ x daily VW consumption (rare!)
- could pile up for very frequently used metadata queries (~10k times)

* *serverless features*

- Snowflake-managed vs user-managed resources
- charged w/ credits, ~WH, depending on the service
- billed per second, not when idle (like for the user-managed resources)
- Snowflake-managed compute + cloud service credits

*Understand Price Tables
for Storage and
Data Transfer*

Tip #25

Checkpoints

* *storage*

- small fraction, in average TB/mo, for compressed data
- On-Demand / Capacity
- Hybrid Tables storage billed separately → GB/mo

* *data transfer*

- very small fraction, in TB transferred, ~the cloud provider
- for data out only (data in free)
- free when within same region/continent/provider
- from COPY, replication, external functions

*Use the
Account Overview Interface
in Snowsight*

Tip #26

Cost Management interface

Account Overview

PREVIEW

Consumption

Budgets

PREVIEW

Resource Monitors

Account spend for **RAA41860** ⓘ from **Feb 24 - Mar 2** ▾

\$1.11

Spend in
USD

0.32

Spend in
credits

\$3

Compute
price/credit

\$0.14

ⓘ Average daily
cost

0.04

Average daily
credits

Consumption

Account Overview PREVIEW **Consumption** Budgets PREVIEW Resource Monitors

🕒 Last 7 days ▾

🛡️ All Accounts ▾

🏷️ All Tags ▾

📄 All Usage Types ▾




\$1.24 USD Spent

By Day ▾

■ RAA41860 ■ XTRACTPRO_BC ■ XTRACTPRO_STD

\$0.6000



Budgets

Account Overview

PREVIEW

Consumption

Budgets

PREVIEW

Resource Monitors

🕒 Current month ▼

ⓘ Data latency can be up to 6 hours



0.2 Credits Used

Mar 1 - Mar 31

Spend

2%

0.2 of 10 Credits

Interval

6%

2 of 31 Days

● Aggregated account spend ● Projected daily spend ● Spending limit



Checkpoints

- * **Cost Management** - Admin interface, w/ date range
 - **Account Overview** - PuPr
 - **Consumption** - per date range/acct, C/S/DT
 - **Budgets** - PuPr, only paid accounts
 - **Resource Monitors** -
- * **Account Overview Insights**
 - *Account budget for current month* - w/ budget(s)
 - *Top warehouses by cost* - w/ SQL
 - *Top databases by storage* - w/ SQL
 - *Most expensive queries* - total/avg exec time + #

*Use
Organization
Accounts*

Tip #27

Accounts interface

Accounts **Replication** PREVIEW

[+ Account](#)

Active Accounts Dropped Accounts

Edition **All**

Cloud **All**

Region **All**

3 Accounts ?



ACCOUNT	EDITION	CLOUD	REGION	CREAT...	LOCATOR	ORGADMIN EN... <small>↑</small>	
RAA41860	Enterprise	AWS	US West (O...	2 years ...	WWA...	✓	...
XTRACTP...	Standard	AWS	US West (O...	2 month...	VXB9...	✓	...
XTRACTP...	Business ...	AWS	US West (O...	2 days a...	DAB7...	✓	...

Accounts interface

- * Create New Account

- *edition*: Standard/Enterprise/Business Critical
- *cloud/region*: AWS/Azure/GCP

- * accounts: active/dropped ← 3..90 days grace period

- manage URLs: locator/org-acct
- enable/disable ORGADMIN

Checkpoints

- * create new related paid Snowflake accounts
- * create Standard account for basic query experiments
- * private share data & native apps within org accounts
- * test new PuPr features available only in paid accounts
- * requires higher level of security than free trial accounts
- * **SHOW ORGANIZATION ACCOUNTS** ← w/ ORGADMIN

Limit

*Warehouse Changes
with Access Control*

Tip #28

Checkpoints

- * **USAGE** ← execute/abort queries (w/ eventual auto-resume)
- * **OPERATE** ← execute/abort queries + view crt/past queries + **suspend/resume**
- * **MODIFY** ← **alter properties (including size)** + assign to resource monitors

- * **MONITOR** ← view current/past queries + usage statistics
- * **APPLYBUDGET** ← add/remove from a budget
- * **OWNERSHIP** ← full control (for a single role)
- * **ALL** ← all privileges but OWNERSHIP

*Adjust Default Values
of Account-Level
Parameters*

Tip #29

Checkpoints

* *time travel* ← cannot be disabled for an account!

- **MIN_DATA_RETENTION_TIME_IN_DAYS** - (def 0), can force TT, does not alter current settings
- **DATA_RETENTION_TIME_IN_DAYS** - (def 1) 0..90 Ent / 0..1 Std, set def for database/schema/permanent tables, but cannot force disable w/ 0

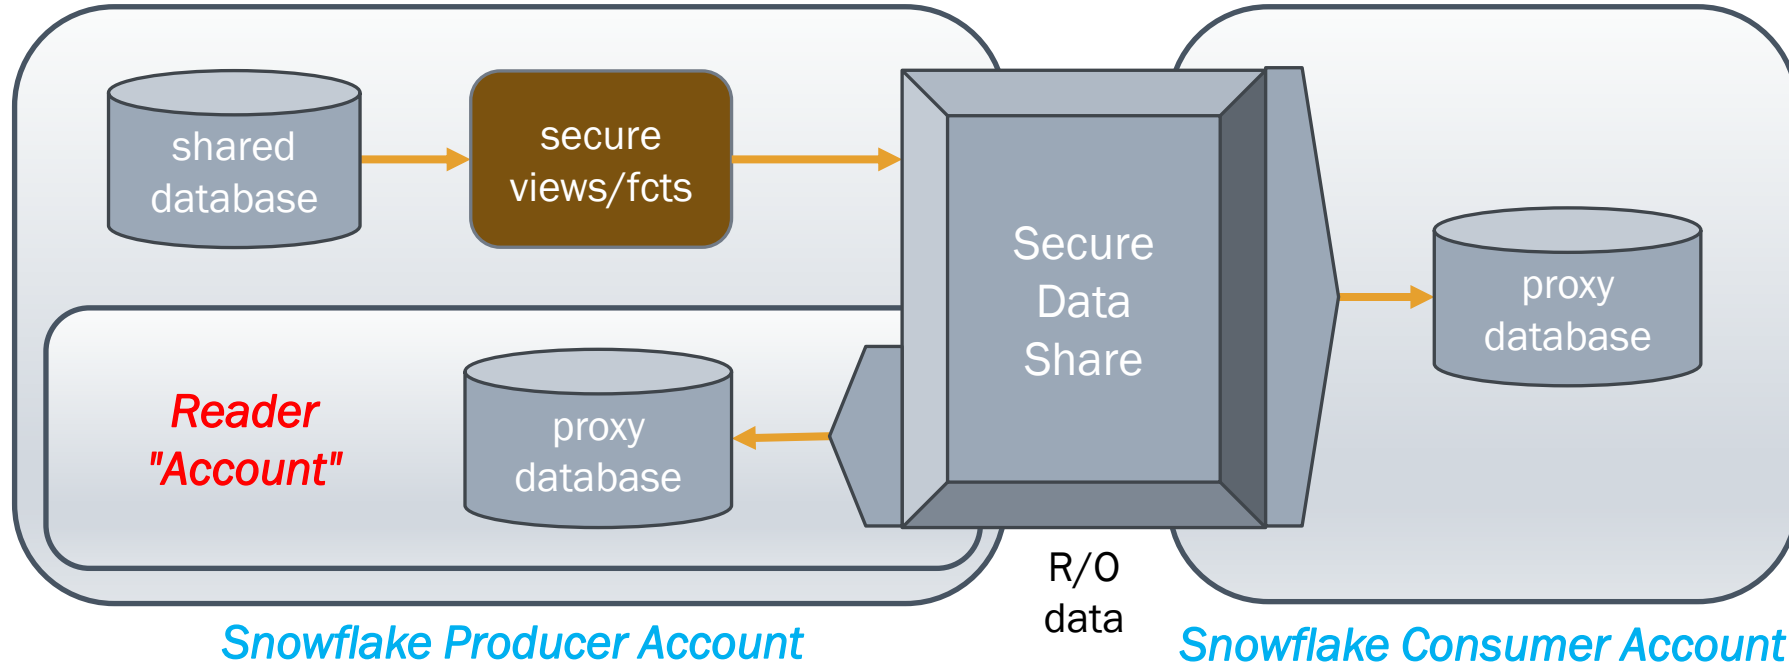
* *replication*

- **INITIAL_REPLICATION_SIZE_LIMIT_IN_TB** – (def 10) max estimated size limit for initial replication of a primary db to a sec db, to prevent accounts w/ sec dbs from accidentally incurring large replication charges

*Careful with
Reader Accounts*

Tip #30

Data Share w/ Reader Account



Checkpoints

- * Reader accounts are not separate Snowflake accounts
- * Reader accounts consume YOUR credits!
- * monitor closely all your Reader accounts
- * set resource monitors w/ budget alerts
- * "local" admins could create/manage WHs there!

Introduction to Snowflake Editions

Tip #31: When to Choose Enterprise over Standard Edition

Tip #32: How to Avoid Multi-Cluster Warehouses

Tip #33: When to Use Incremental Materializations

Tip #34: How to Emulate Materialized Views

Tip #35: The Case for Extended Time Travel

Tip #36: Use Standard Edition Account for Analytics

Tip #37: Use Separate Standard Edition Account for Common Queries

Tip #38: How to Reduce Costs to Zero for an Inactive Paid Account

Tip #39: When to Choose the Business Critical Edition

Tip #40: When to Choose the Virtual Private Snowflake (VPS) Edition

*When to Choose
Enterprise over
Standard Edition*

Tip #31

Essential New Features

- * **Multi-Cluster Virtual Warehouses** → w/ auto-scaling out for concurrency
- * **Query Acceleration Service** → boost WH size Nx, for eligible queries
- * **Search Optimization Service** → for point lookup queries
- * **Materialized Views** → w/ auto-maintenance of results
- * **Extended Time Travel** → max 90 days (from 1) for permanent tables

Other New Features

- * [Data Masking Policies](#) → column-level security in tables/views
- * [Row Access Policies](#) → row-level security for query rows
- * [Object Tagging](#) → to track sensitive data + resource usage
- * [Classification](#) → to auto-tag potentially sensitive data
- * [User Access Auditing](#) → in ACCESS_HISTORY view
- * [Periodic Rekeying of Encrypted Data](#) → increased protection
- * [Early Access to Weekly New Releases](#) → 12-hour

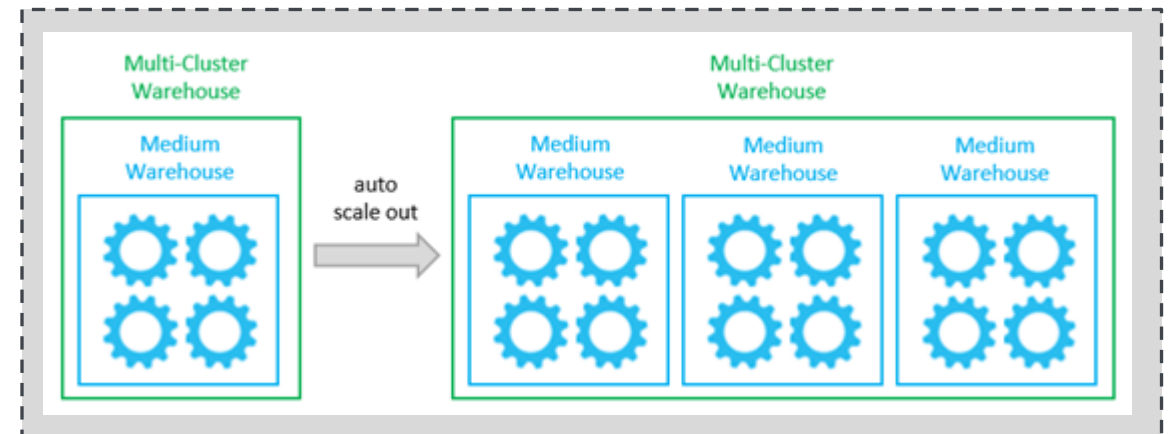
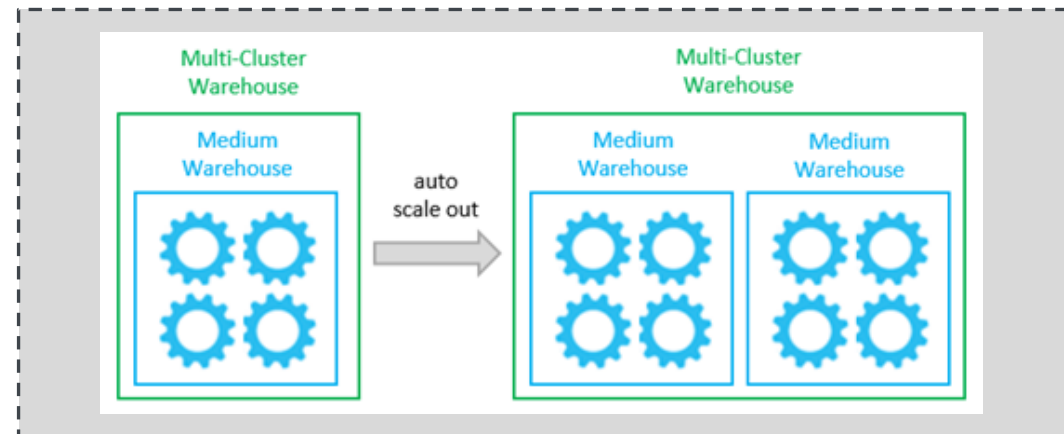
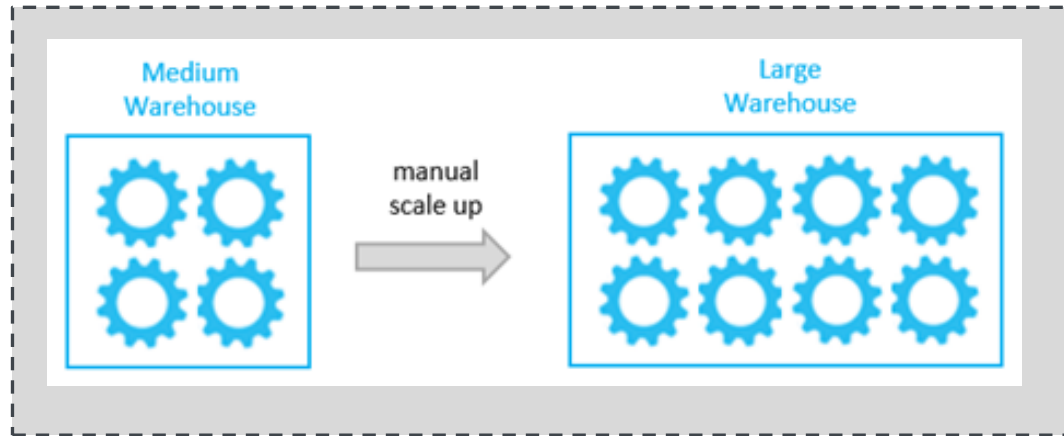
Checkpoints

- * SMB → large enterprise/org
- * 1.5x credit cost
- * higher number of users → multi-cluster WHs (concurrency)
- * QAS + SOS + MVs → improved query performance
- * 90 days TT → better data protection
- * column/row policies, obj tagging, classif → better data gov
- * rekey, user audit → better security

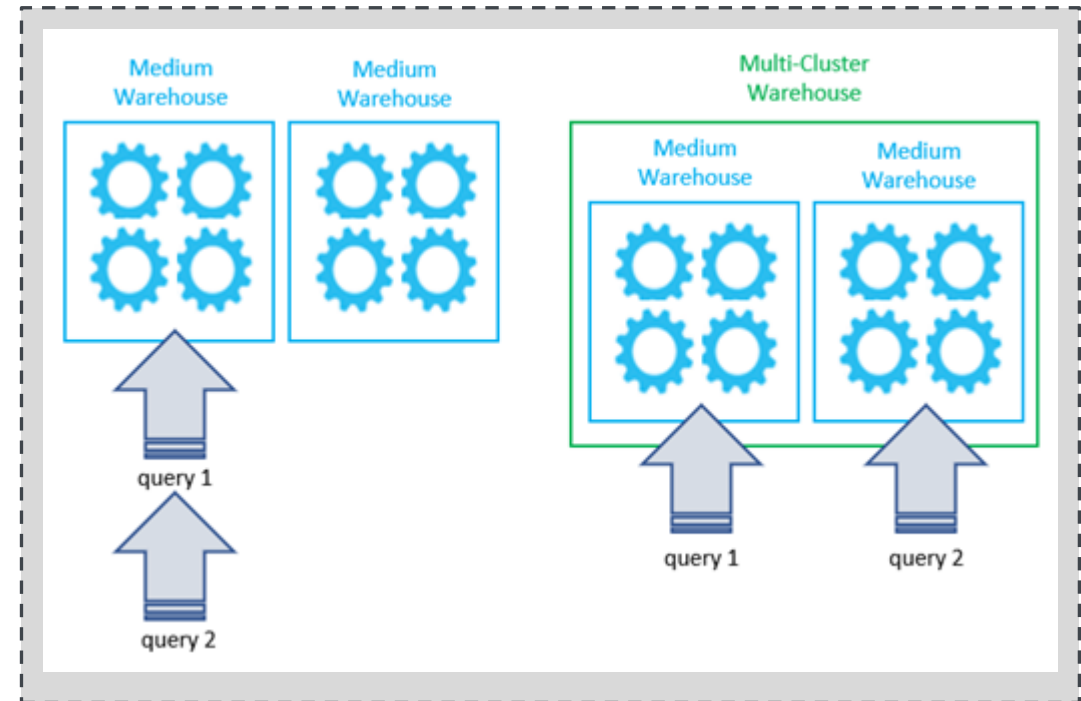

*How to Avoid
Multi-Cluster
Warehouses*

Tip #32


Scalability



Concurrency



Custom Auto-Router



Checkpoints

- * multi-cluster WHs (Enterprise+ only) improve concurrent access (avoid queuing)
- * *custom auto-router* = ~ load balancer, on standard WHs
- * auto-scale out/in to standard overflow WHs (manual scale up/down normally)
- * could reroute if M+ queries on current WH (use MAX_CONCURRENCY_LEVEL)
- * auto-route work on load spikes (to existing/dynamic WHs w/ same/different sizes)
- * reroute to pool of WHs (10+!) after STATEMENT_QUEUED_TIMEOUT_IN_SECONDS expires
- * could use default auto-suspend times (multi-cluster could use longer time!)

*When to Use
Incremental
Materializations*

Tip #33

Materialized Views

- * Enterprise+, storage + compute costs (10 credits/h) → when refreshed
- * act more like tables, as they store R/O data → result never outdated
- * base table data changes → back incremental update
- * based on one single table (self-joins not allowed)
- * can be secured (not optimized) + can have a (different) cluster key
- * has limited SQL syntax, no UDFs

Checkpoints

- * when query result doesn't change often (base table not freq updated)
- * when results are asked often (less complex query executed here)
- * when original query consumes lots of resources
- * to improve performance on an external query
- * when lots of aggregations returned
- * when few rows constantly returned only from a very large table
- * need a different cluster key than the base table

How to Emulate Materialized Views

Tip #34

Checkpoints

- * use a table instead, instantly/periodically updated
- * use a task to periodically update the table
- * task + CHANGE_TRACKING enabled on the source table
- * task + STREAM ← legacy mechanism for CDC
- * use a Snowflake *dynamic table* instead ← new!

*The Case for
Extended Time Travel*

Tip #35

Time Travel Use Cases

* *looking back in time*

- SELECT ... FROM ... AT (TIMESTAMP => timestamp) ...
- SELECT ... FROM ... AT (OFFSET => rel_time_diff) ...
- SELECT ... FROM ... AT (STREAM => 'name') ...
- SELECT ... FROM ... AT/BEFORE (STATEMENT => query_id) ...

* *in zero-copy cloning*

- CREATE DATABASE/SCHEMA/TABLE target CLONE source AT/BEFORE (...)

* *restoring dropped objects*

- DROP DATABASE/SCHEMA/TABLE name
- SHOW DATABASES/SCHEMAS/TABLES HISTORY [...] ← dropped
- UNDROP DATABASE/SCHEMA/TABLE name ← restore dropped object

Checkpoints

- * *time travel*

- max 90 days for permanent tables in Enterprise Edition
- max 1 day for permanent tables in Standard Edition
- max 1 day for transient/temporary tables (in any edition)

- * for DATABASE → SCHEMA → TABLE


- * `DATA_RETENTION_TIME_IN_DAYS` → in CREATE ... / ALTER ... SET ...

- set to 0 to disable → but cannot permanently disable in account!

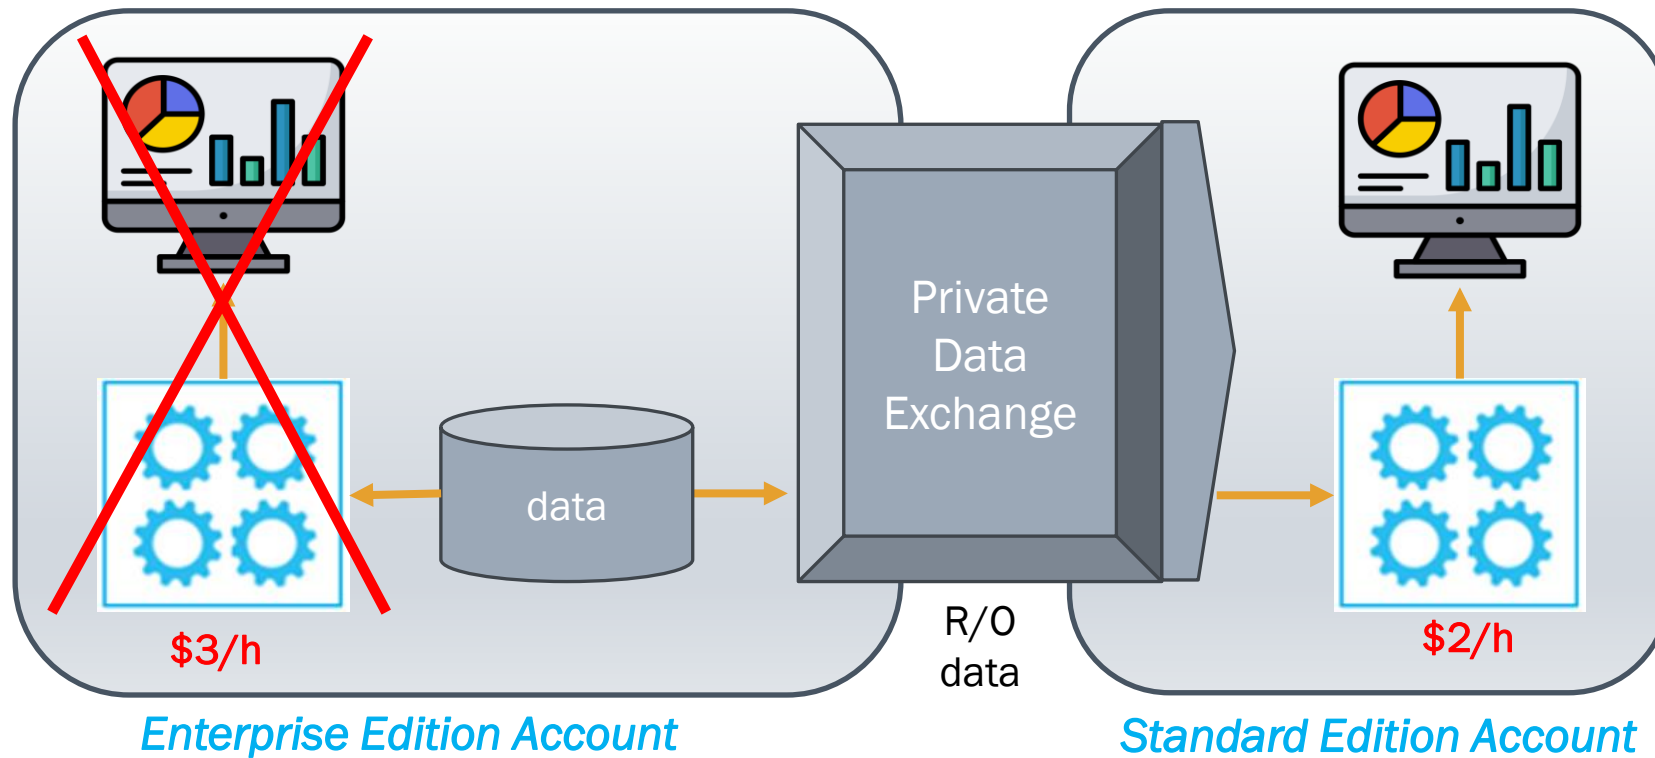
*Use Standard Edition
Account for Analytics*

Tip #36

Analytics in Enterprise Edition



Analytics in Standard Edition



Checkpoints

- * w/ [private data share](#) Enterprise → Standard accounts
- * analytics in Standard Edition consume just \$2/h credits
- * save 30% on compute costs!
- * *limitations*: multi-cluster WHs, mat views, QA, SO, clustering
- * multi-cluster warehouses could be avoided
- * keep data w/ mat views + clustering in Enterprise

*Use Separate Standard
Edition Account
for Common Queries*

Tip #37

Checkpoints

- * any Organization must also have one Standard account
- * for higher edition account → create also one Standard
- * if no special features needed → run query on Standard
- * *credit cost*: \$2 Standard, \$3 Enterprise, \$4 Business
- * no need for a multi-cluster warehouse if you're alone

*How to Reduce
Costs to Zero for
an Inactive Paid Account*

Tip #38

Checkpoints

- * do not share it with other people
- * keep one single 1-Node X-Small WH w/ auto-suspend 1 minute
- * associate your user connection w/ this WH
- * do not run long queries (Cancel them if they already take too long)
- * do not have active/scheduled tasks running (better drop them all)
- * do not leave any serverless features running

Checkpoints (2)

- * drop any potential dangerous object after experiments, even if inactive (WHs, tasks)
- * don't check too much on metadata (in screen or by SQL)
- * remove any app installed from the Marketplace after testing
- * create explicit transitory tables w/ no time travel
- * change default days for time travel to 0 at the account level
- * do not store/copy/keep/transfer large data

*When to Choose the
Business Critical Edition*

Tip #39

Enhanced Data Protection

- * Failover/failback Snowflake accounts → w/ redir conns
- * Tri-Secret Secure w/ CMKs
- * AWS PrivateLink support
- * Ext Funcs through AWS API Gateway Private Endpoints

Support for Regulations

- * PHI Data Support (for HIPAA and HITRUST CSF)
 - for orgs w/ extremely sensitive data, as in healthcare
 - required signed business agreement w/ Snowflake
- * PCI DSS Compliance (payment card industry)
- * Support for public sector (FedRAMP and ITAR)
- * Support for IRAP - Protected (P) data (in Asia Pacific)

Checkpoints

- * former ESD (Enterprise for Sensitive Data)
- * Enterprise+, 2x compute credit cost over Standard
- * enhanced security and data protection
- * failover/failback between Snowflake accounts
- * for business continuity/disaster recovery
- * support for regulations (PHI/HIPAA, PCI, US Federal)
- * combine w/ Standard/Enterprise for reg queries → 50% cost savings!

*When to Choose the
Virtual Private Snowflake
(VPS) Edition*

Tip #40

New Features

- * completely separate Snowflake environment
- * isolated from all other Snowflake accounts
- * dedicated virtual servers (w/ in-memory encryption key)
- * pool of compute resources (used in VWs)
- * dedicated metadata store
- * secure data share (w/ Marketplace/Data Exchange) disabled by def
- * call support to enable data sharing with non-VPS customers

Credit Cost

- * Standard: \$2.00 - \$3.10
- * Enterprise: \$3.00 - \$4.65
- * Business Critical: \$4.00 - \$6.20
- * VPS: \$6.00 - \$9.30

Checkpoints

- * Business Critical+, 2x credit cost of Enterprise (\$6+ in AWS/US)
- * completely isolated/separated account/env → highest level of security
- * dedicated servers, metadata, pool of VW resources
- * secure data share disabled by def, but can enable
- * 24-hour early access to weekly new releases
- * for orgs w/ strictest reqs: financial inst., large ent. w/ highly sensitive data

Introduction to Query Monitoring

Tip #41: Monitor Longest Running Queries

Tip #42: Interpret Query History

Tip #43: More Charts for Query Monitoring

Tip #44: Use Query Tags

Tip #45: Reduce Frequency of Simple Queries

Tip #46: Reduce Frequency of Metadata Queries

Tip #47: Reduce Frequency of SHOW Commands

Tip #48: Clone Less Frequently

Tip #49: Change Query Schedules

Tip #50: Parallel over Sequential Transfer and Processing

*Monitor
Longest Running
Queries*

Tip #41

Checkpoints

- * from `snowflake.account_usage.QUERY_HISTORY` view
- * top N w/ longest individual *execution_time*
- * top N w/ longest AVG exec time, by *type/WH size/user/...*
- * top N query types w/ most consumed *credits*
- * top N w/ longest *total* execution time overall
- * top N most *frequently executed* queries

*Interpret
Query History*

Tip #42

Checkpoints

- * Monitoring - [Query History](#) ← in Snowsights (14 days)
- * ACCOUNT_USAGE.[QUERY_HISTORY](#) view (1 year)
 - *total_elapsed_time* = compilation + execution + queued...
 - *queued_provisioning/repair/overload_time*
 - *transaction_blocked_time*
 - *transferred bytes* = ...
- * INFORMATION_SCHEMA.[QUERY_HISTORY\[_BY...\]](#) functions (7 days)
 - filter by session/user/warehouse
 - less info

*More Charts
for Query Monitoring*

Tip #43

Checkpoints

- * histogram of queries duration - in seconds + log scale
- * zoom into top-3 longest queries in detail
- * time-histograms of aggregate queries duration - in seconds
- * day-hour histogram
- * longest and most frequent queries - with log scales

Use
Query Tags

Tip #44

Checkpoints

- * **QUERY_TAG** = session-level param, 2k chars max, can use JSON
 - `ALTER SESSION SET QUERY_TAG = '...'` ← w/ SQL
 - `session_parameters={'QUERY_TAG': '...'}` ← w/ Python Connector
 - `session.query_tag = '...'` ← w/ Snowpark
- * filter QUERY_TAG column in QUERY_HISTORY
- * alternative: end query w/ '-- comment'
- * not a session variable (no SET QUERY_TAG = ...)
- * no object tagging (no CREATE TAG ...)

*Reduce Frequency
of Simple Queries*

Tip #45

Checkpoints

- * ~10K calls/day → significant cloud service usage
- * for `SELECT 1 / seq1.NEXTVAL / CURRENT_SESSION()`
- * session ID does not change while a connection is open
- * could save `CURRENT_SESSION()` value after first call
- * JDBC SnowflakeConnection.***getSessionId()*** uses cache!
- * could use a dashboard chart, to monitor

*Reduce Frequency
of Metadata Queries*

Tip #46

Checkpoints

- * ~10K calls/day → significant cloud service usage
- * query INFORMATION_SCHEMA → uses cloud services
- * query ACCOUNT_USAGE → uses a virtual warehouse
- * ACCOUNT_USAGE has more data → 1y, for account
- * could use a dashboard chart, to monitor

*Reduce Frequency
of SHOW Commands*

Tip #47

SHOW for RBAC

SHOW ROLES → builtin+custom roles

SHOW GRANTS TO ROLE ... → assigned roles+privs

SHOW USERS → users

SHOW GRANTS TO USER ... → assigned roles

SHOW for ERDs

SHOW TABLES → tables

SHOW COLUMNS → table columns

SHOW PRIMARY KEYS → PRIMARY KEY constraints

SHOW UNIQUE KEYS → UNIQUE constraints

SHOW IMPORTED KEYS → FOREIGN KEY constraints

Checkpoints

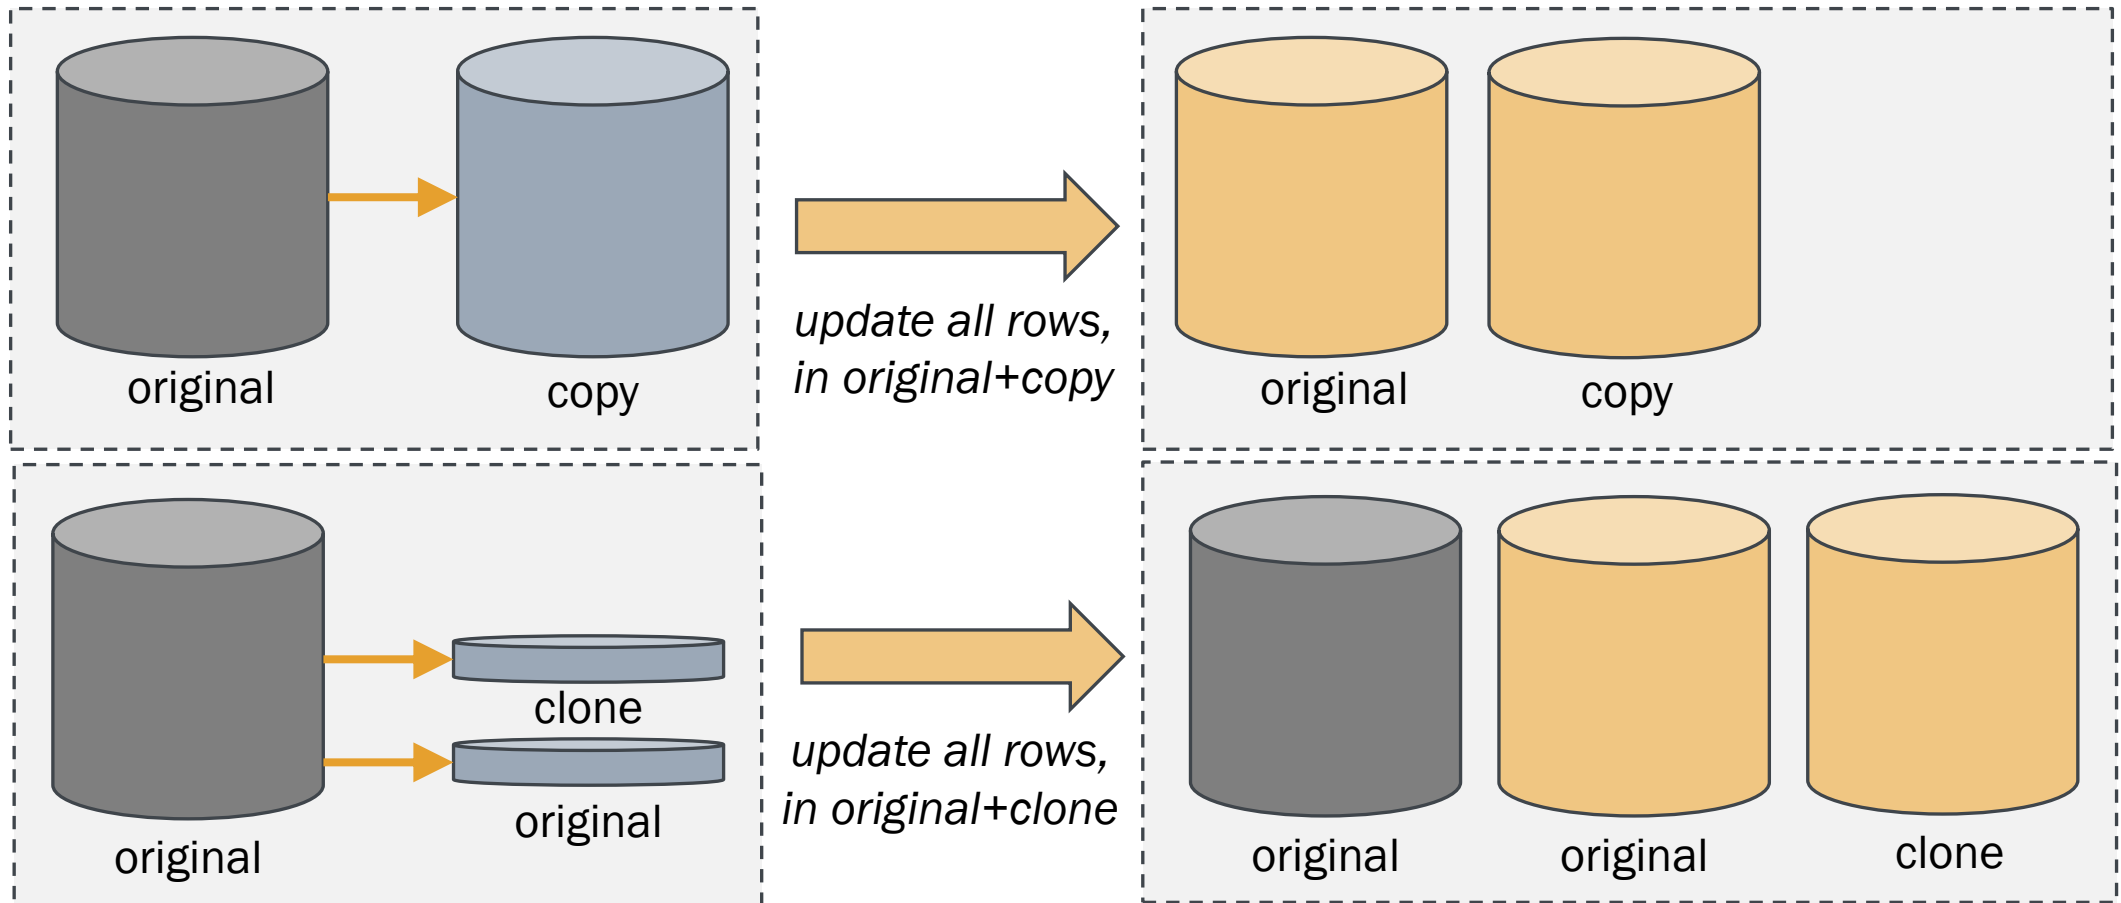
- * SHOW cmds → consume only cloud service resources
- * avoid SHOW cmds at a high frequency
- * check third-party tools, on top of Snowflake
- * check partner tools
- * could use a dashboard chart, to monitor

Clone

Less Frequently

Tip #48

Changing Most Rows in Common



Checkpoints

- * extensive DDL (cloning) → significant cloud serv usage
- * DDL ops (like cloning) use only cloud compute serv
- * do not frequently create/drop large schemas/tables
- * do not frequently clone databases for backup
- * avoid cloning if you intend to change most rows in common
- * updating most rows in common will increase total storage space

*Change
Query Schedules*

Tip #49

Cost

* *task*

- scheduled every 2 mins, runs for max 2 secs each time
- $(24\text{h} \times 60 \text{ mins} / 2 \text{ mins}) \times 2 \text{ secs} = 1440 \text{ secs} = 24\text{min/day}$
- 24min/day for a related WH $\rightarrow 0.5 \text{ credits/day} = \$1.5/\text{day}$

* *warehouse*

- X-Small, 1 credit/h (\$3/credit), auto-suspended after 1 min inactivity
- resumed by the task every 2 mins = 30mins/h = 12h/day
- $24\text{h} \times 30 \text{ mins} = 720\text{min} = 12\text{h} \rightarrow \$36/\text{day}$

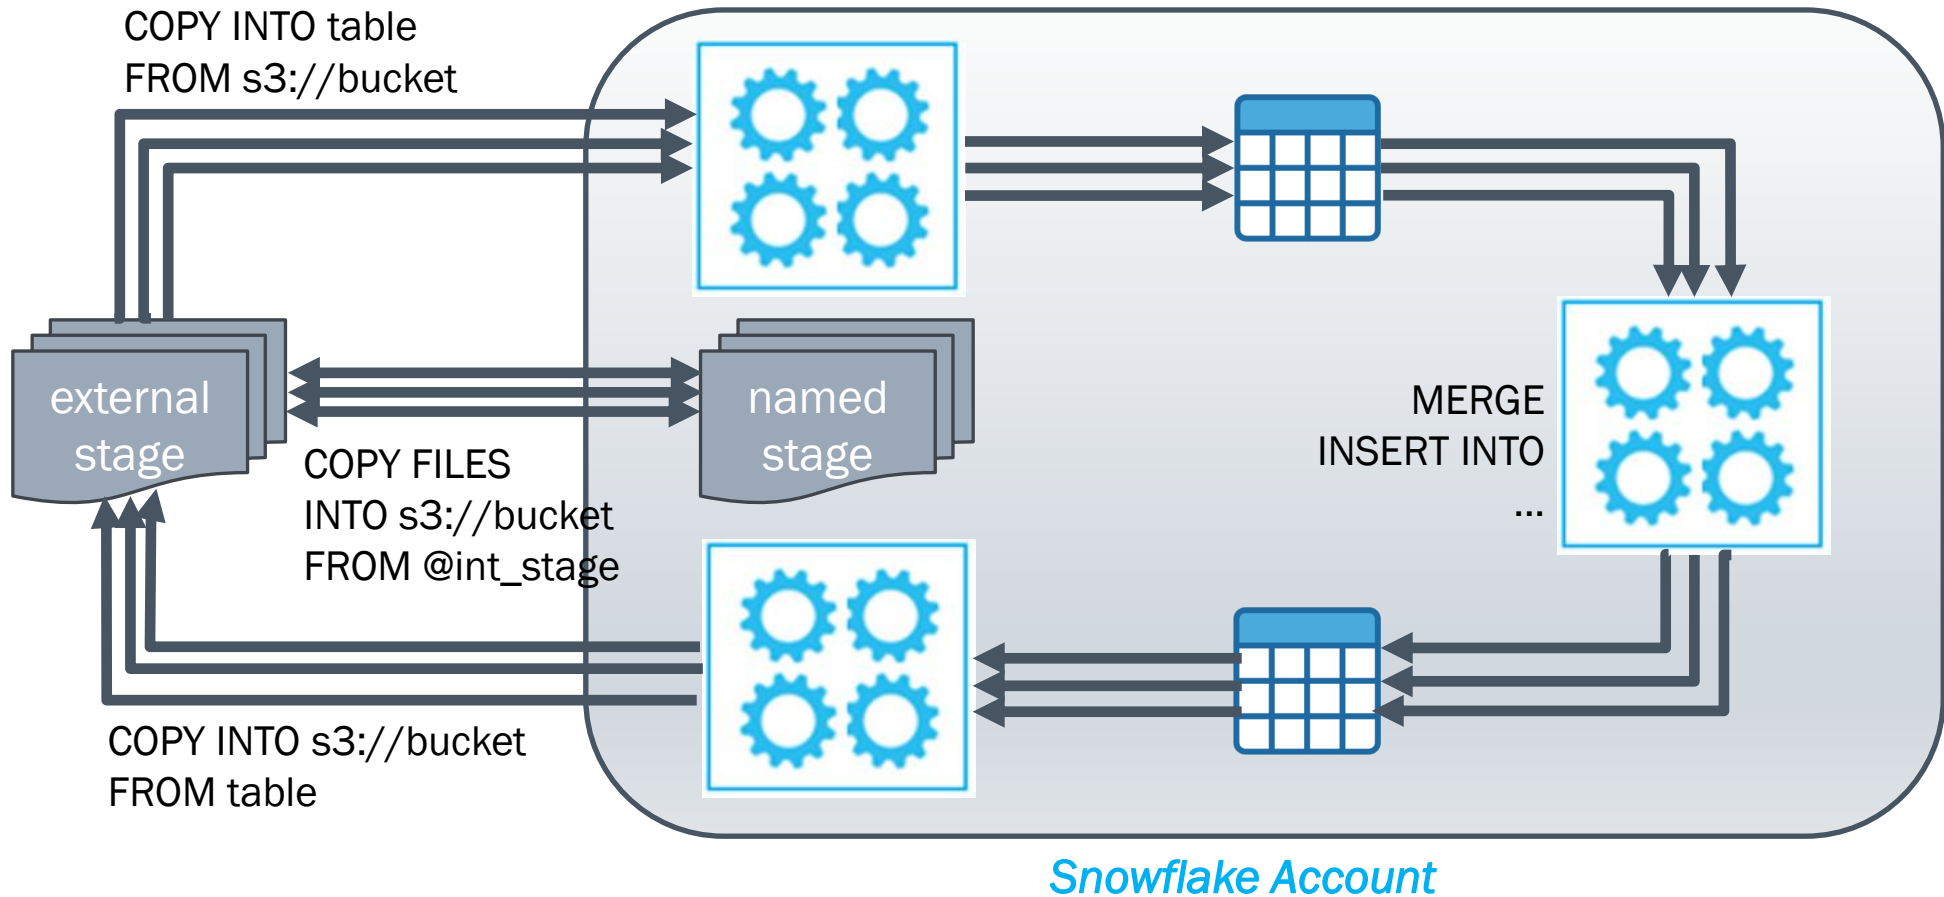
Checkpoints

- * frequently scheduled tasks not ok w/ managed WHs
- * you may pay more for the WH idle, suspending (1 min)
- * *solution*: try using a WH already up, w/ other queries
- * adjust eventually the schedule, to find such a WH up
- * avoid up-down pattern → exec more queries together
- * consider a serverless task (no WH, but 1.x the price)

*Parallel over Sequential
Transfer and Processing*

Tip #50

Parallel Processing



Checkpoints

- * use parallel query processing → max_concurrency_level at WH
- * COPY INTO table from multiple split input files
- * COPY INTO @stage w/ multiple split output files
- * COPY FILES w/ multiple files, from stage to stage
- * MERGE over individual INSERT/UPDATE/DELETE
- * INSERT INTO multiple tables w/ parallel processing
- * batch transformations w/ multiple SQL connections

Introduction to Query Optimization

Tip #51: Use the Query Profile

Tip #52: Use the Explain Statement

Tip #53: Use Data Caching

Tip #54: Queries on Data Lakes

Tip #55: Use Vectorized Python UDFs

Tip #56: Use Batch Commands to Prevent Transaction Locks

Tip #57: Reduce Query Complexity and Compilation Time

Tip #58: Check for Cross Joins and Exploding Joins

Tip #59: Process Only New or Updated Data

Tip #60: Remote Spillage Optimization

*Use the
Query Profile*

Tip #51

Checkpoints

- * generates graph of the actual execution plan
- * can use for running queries as well
- * based on the generated query ID
- * carefull w/ queries executed from cache
- * great for TableScan and Join steps ← most costly
- * evaluate partitions scanned/total + bytes scanned
- * check number of rows returned by a Join → exploding joins?

*Use the
Explain
Statement*

Tip #52


Checkpoints

- * get the logical execution plan without running the query
- * estimates micro-partitions to scan and total bytes
- * USING TEXT/JSON → hierarchy of logical steps
- * TableScan steps could use lots of data
- * consider columns used for Automatic Clustering
- * see when mat views could be scanned in the back
- * could use as query syntax check

Use
Data Caching

Tip #53

Query Result Caching



Checkpoints

HOT

- result cache on
- cache hit → `SELECT * FROM TABLE(RESULT_SCAN(LAST_QUERY_ID()))`

WARM

- result cache off or no cache hit
- warehouse up and result still in the warehouse (in RAM or local disks)
- get query result from warehouse cache (`BYTES_SCANNED > 0`)

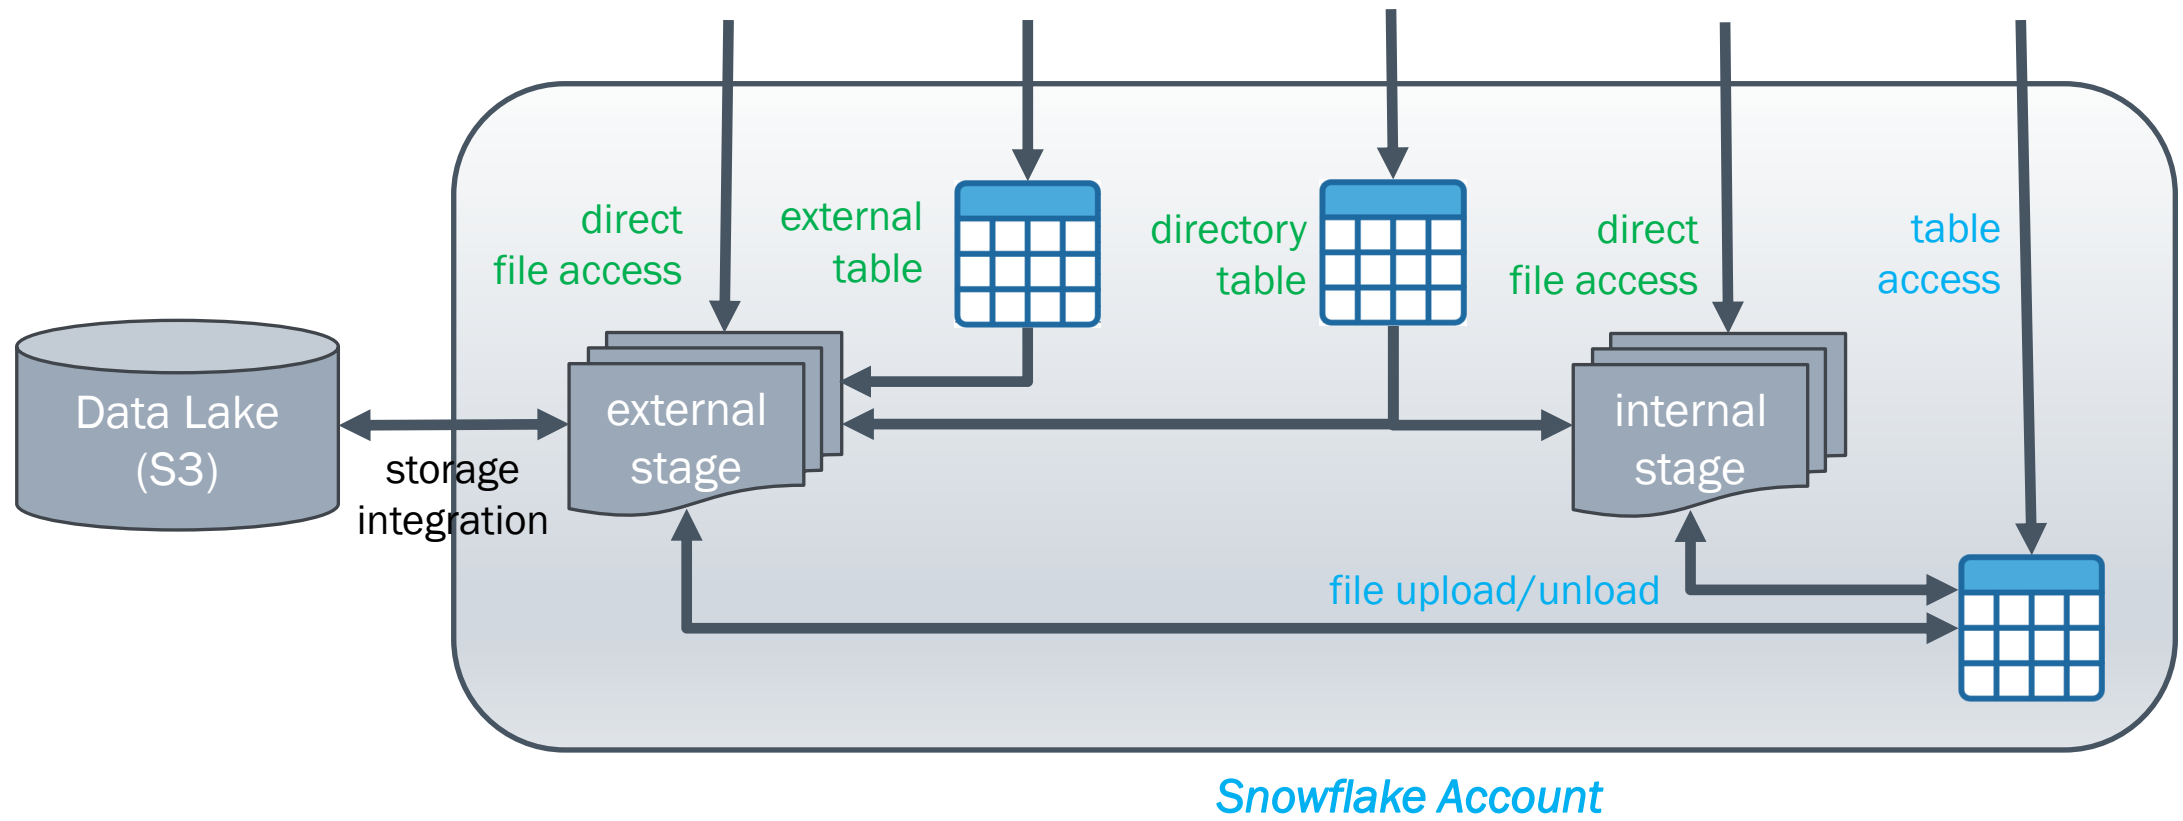
COLD

- no warehouse cache (warehouse suspended, not up)
- result cache off (`ALTER SESSION SET USE_CACHED_RESULT = false`) or no cache hit
- must run query (could also be a new query)

*Queries on
Data Lakes*

Tip #54

Data Warehouse vs Data Lake



Checkpoints

- * could be better to leave files in stage and query from there
- * *data warehouse*
 - *file upload/unload* → between stage files and tables (partitions!)
 - *table access* → typical SQL on stored Snowflake data
- * *data lake*
 - *direct file access* → on internal/external stage files (\$1, \$2...)
 - JSON in VARIANT is more optimized!
 - Parquet vs CSV: 50x+ times faster!
 - *external tables* → on external stage (w/ column defs!)
 - *directory tables* → for unstructured data (refs only in the database!)

*Use Vectorized
Python UDFs*

Tip #55

Checkpoints

- * row-by-row processing → batches of rows (better performance!)
- * **non-vectorized UDF** = scalars → UDF → scalar return
- * **vectorized UDF** = scalars → **pandas DataFrames** → UDF → **pandas arrays/Series** → scalar returns

```
@vectorized(input=pandas.DataFrame)
```

optional max_batch_size

*Use Batch Commands
to Prevent
Transaction Locks*

Tip #56

Checkpoints

- * UPDATE/MERGE → transaction lock on a table
 - other cmds cannot exec on table until lock released
 - queries consume cloud serv credits waiting for lock
- * *solution*: use a batch cmd instead of single updates
 - use multi-row batch INSERT in temporary stage table
 - use task to periodically UPDATE/MERGE into target table

*Reduce Query Complexity
and Compilation Time*

Tip #57

Cost

* *object dependencies query*

- compilation: 19 secs
- execution: 5.6 secs
- rows returned: 0

* *data lineage query*

- compilation: 1.2 secs
- execution: 4.2 secs
- rows returned: 1.4K

Checkpoints

* *complex queries*

- can consume significant cloud services compute resources
- will have high compilation times
- see *compilation_time* vs *execution_time* in QUERY_HISTORY

* *examples*

- queries with lots of joins
- queries with cartesian products
- queries with IN operator and large lists
- very large queries

*Check for Cross Joins
and Exploding Joins*

Tip #58

Cross Joins vs Exploding Joins

C_CUSTKEY	C_NAME	C_NATIONKEY	C_ACCTBAL
30001	Customer#000030001	4	8848.47
30002	Customer#000030002	11	5221.81
30003	Customer#000030003	21	3014.89
30004	Customer#000030004	23	3308.55
30007	Customer#000030007	5	3912.67
30010	Customer#000030010	21	8599.71
30011	Customer#000030011	1	4442.02
30012	Customer#000030012	3	9027.69
30013	Customer#000030013	4	3438.09



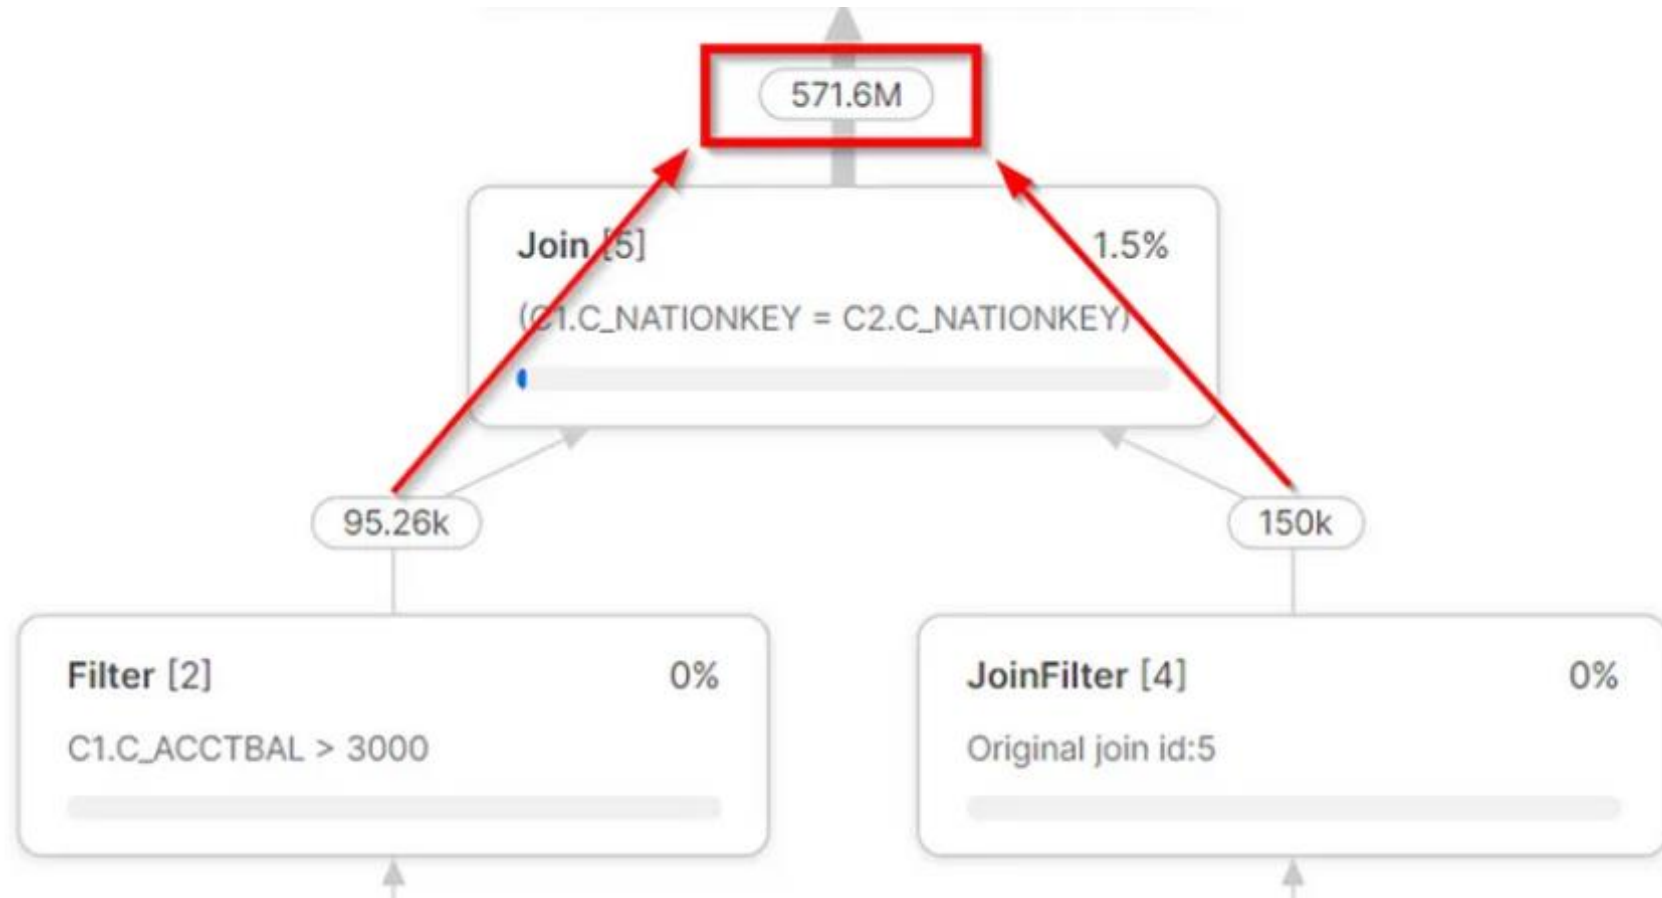
C_CUSTKEY	C_NAME	C_NATIONKEY	C_ACCTBAL
30001	Customer#000030001	4	8848.47
30002	Customer#000030002	11	5221.81
30003	Customer#000030003	21	3014.89
30004	Customer#000030004	23	3308.55
30007	Customer#000030007	5	3912.67
30010	Customer#000030010	21	8599.71
30011	Customer#000030011	1	4442.02
30012	Customer#000030012	3	9027.69
30013	Customer#000030013	4	3438.09

C_CUSTKEY	C_NAME	C_NATIONKEY	C_ACCTBAL
30001	Customer#000030001	4	8848.47
30002	Customer#000030002	11	5221.81
30003	Customer#000030003	21	3014.89
30004	Customer#000030004	23	3308.55
30007	Customer#000030007	5	3912.67
30010	Customer#000030010	21	8599.71
30011	Customer#000030011	1	4442.02
30012	Customer#000030012	3	9027.69
30013	Customer#000030013	4	3438.09



C_CUSTKEY	C_NAME	C_NATIONKEY	C_ACCTBAL
30001	Customer#000030001	4	8848.47
30002	Customer#000030002	11	5221.81
30003	Customer#000030003	21	3014.89
30004	Customer#000030004	23	3308.55
30007	Customer#000030007	5	3912.67
30010	Customer#000030010	21	8599.71
30011	Customer#000030011	1	4442.02
30012	Customer#000030012	3	9027.69
30013	Customer#000030013	4	3438.09

Exploding Joins in Query Profile



Joins

- * *one-to-one joins*: 95.26K rows
- * *cross joins*: $95.26\text{K} * 150\text{K} = 9\text{B+ rows}$ (cancelled)
- * *exploding joins*: $< 95.26\text{K} * 150\text{K} = 571.6\text{M rows}$

customer → 150K rows

customer w/ *c_acctbal* > 3000 → 95.26K rows

Checkpoints

- * **exploding joins** = inner join w/ huge number of matches
- * ~cross joins/cartesian product → large number of rows
- * returns close to **N x M rows** (N, M = rows of tables)
- * more related to many-to-many relationships
- * detect by looking at the Query Profile

Process Only

New or Updated Data

Tip #59


Checkpoints

- * only the changed data gets processed → reduce compute costs
- * avoid transferring the whole batch of data, daily → when most data did not change
- * CDC (Change Data Capture) → "play" transaction log for new changes only
- * *incremental processing*: do not rerun entire batches of data for minor updates
- * restart interrupted pipelines from point of interruption → no redundancies
- * on failed data uploading → keep successful loaded data, continue w/ failed
- * on failed batch data transformations → partition and replay for failed ranges only

Remote Spillage Optimization

Tip #60

Spilling to Local/Remote Storage



Cost

- * Large (8 Nodes): 2 mins, 60 GB spilled to local
- * Snowpark-optimized Large (16x memory): 1:30m, ~3 GB spilled to local

on 48.6 GB table

Checkpoints

- * **Bytes spilled to local storage** = RAM → SSD (all in WH!)
- * **Bytes spilled to remote storage** = SSD → blob storage (slowest!)
- * on SSD = intermediate results for ORDER BY, GROUP BY, PIVOT...
- * RAM (fastest + most expensive) >> SSD (disk) >> blob storage (network)
- * *solution*
 - use a bigger WH (= 2x RAM+SSD!) - Snowpark-optimized (16x RAM) or QAS as alternative
 - use less data
- * check **bytes_spilled_to_local/remote_storage** in QUERY_HISTORY

Introduction to Serverless Features

Tip #61: Monitor the Cost of Automated Jobs

Tip #62: Estimate Cost of Scheduled Tasks

Tip #63: When to Use Serverless Tasks

Tip #64: Replace Snowpipe with Snowpipe Streaming

Tip #65: Estimate Cost of Automatic Clustering on Tables

Tip #66: Estimate Cost of the Query Acceleration Service (QAS)

Tip #67: Estimate Cost of the Search Optimization Service (SOS)

Tip #68: Reduce Materialized Views Maintenance Cost

Tip #69: Reduce Database Replication Cost

Tip #70: Estimate Cost of Hybrid Tables

*Monitor the Cost
of Automated Jobs*

Tip #61

Cost

Table 5: Serverless Feature Credit Table		
Feature	Snowflake Credits per Compute-Hour	
	Snowflake-managed compute	Cloud Services
Clustered tables	2	1
Copy Files ¹⁰	2	N/A
Hybrid Tables requests ¹⁰	1	1 In addition: 1 Credit per 30GB read 1 Credit per 7.5GB write
Logging	1.25	1
Materialized views maintenance	10	5
Materialized views maintenance in secondary databases	2	1
Query acceleration	1	1
Replication	2	1
Search optimization service	10	5
Search optimization service in secondary databases	2	1
Serverless tasks	1.2	1
Snowpipe	1.25	N/A but charged 0.06 Snowflake Credits/1000 Files
Snowpipe Streaming	1	N/A but charged at an hourly rate of 0.01 Snowflake Credits per client instance

METERING_HISTORY View

- * *start_time* .. *end_time* = hourly credit usage, for 1y
- * *service_type* + *name*
- * *credits_used_compute* = WH + serverless
- * *credits_used_cloud_services* = ~1 credit/h (10 for MVs/SO)
- * *credits_used* = *credits_used_compute* + *credits_used_cloud_services*
- * *bytes* + *rows* + *files* – for auto-clustering / pipe / snowpipe_streaming

Service Types

- * [WAREHOUSE_METERING\[_READER\]](#) — any edition, w/ WH name [+for Reader accounts]
- * [SERVERLESS_TASK](#) — not user-managed, ~1.5x more expensive (1.2x now)
- * [AUTO_CLUSTERING](#) — any edition, auto-background maintenance of clustered table/MV (+initial clustering+reclustering), ~2 credits/h
- * [MATERIALIZED_VIEW](#) — 10 credits/h [2 in sec dbs] + 5 for cloud → storage+compute (per second). when base table changes → MVs updated by serverless back service
- * [QUERY_ACCELERATION](#) — scale factor to limit cost → N nodes x boost
- * [SEARCH_OPTIMIZATION](#) — 10 credits/h [2 in sec dbs] + 5 for cloud → storage+compute maintenance (higher cost when large table data changes + deletes + auto-clustering)

Service Types (cont.)

- * **PIPE + SNOWPIPE_STREAMING** — Snowpipe ~1.25 credits/h + 0.06 credits/1000 files. Streaming w/ 2 table+service entries, 1 credit/h + 0.01 credits/h/client.
- * **HYBRID_TABLE_REQUESTS** — 1 credit/h per req + 1 cloud credit per 30GB R/7.5GB W
- * **COPY_FILES** — for COPY INTO between stages, 2 credits/h (no cloud)
- * **REPLICATION** — for dbs/shares (any edition) / failover across accounts (BC) – 2 credits/h
- * **SNOWPARK_CONTAINERS** — SPCS, all editions → costs for storage (image repository+logs+mounting volumes) + compute pool (1+ VM nodes) + data transfer (outbound + cross-AZ)
- * **AI_SERVICES** — compute cost for Cortex LLM functions → per number of tokens (~4 chars) processed by the calls

Checkpoints

- * automated jobs → ~70% of all compute cost
- * credits per compute-hour
- * credits used = Snowflake-managed compute + cloud services
- * *cloud services*: for AuthN / infra/metadata management / query parsing+optimization / access control...
- * `METERING_HISTORY` view

Estimate Cost of Scheduled Tasks

Tip #62

Cost

* *managed task*: 1 sec (exec time) x 60 x 24 = <30 min → <**\$1.50**/day (estimated)

* *warehouse*: 1 min (up every time) x 60 x 24 = 24 h → \$3 x 24h = **\$72**/day

user-managed task w/ 1 sec total exec time + executed every minute

X-Small dedicated WH w/ 1 min auto-suspend, at \$3/h/node credit (in AWS/US)

Checkpoints

- * *user-managed* (w/ WH) vs *serverless* tasks (no WH)
 - `USER_TASK_MANAGED_INITIAL_WAREHOUSE_SIZE` → **def MEDIUM!**
- * **WHEN** sql_condition → just cloud service credits!
 - task not triggered at all if WHEN not satisfied
- * **SCHEDULE** → not for child tasks, w/ manual EXECUTE TASK if none
- * `USER_TASK_TIMEOUT_MS` → **def 1h** (max 1 day)
 - `STATEMENT_TIMEOUT_IN_SECONDS` → sometimes w/ higher priority, in WH
- * `information_schema.TASK_HISTORY()`
- * check orphan tasks running but not being used

When to Use Serverless Tasks

Tip #63

Cost

- * *serverless task*: 1 sec (exec time) x 60 x 24 = <30 min → <**\$1.80**/day
- * *managed task*: 1 sec (exec time) x 60 x 24 = <30 min → <**\$1.50**/day (estimated)
- * *warehouse*: 1 min (up every time) x 60 x 24 = 24 h → \$3 x 24h = **\$72**/day

serverless task w/ 1 sec total exec time + executed every minute → \$3.6/h/node

user-managed task w/ 1 sec total exec time + executed every minute

X-Small dedicated WH w/ 1 min auto-suspend, at \$3/h/node credit (in AWS/US)


Checkpoints

- * favor serverless tasks for *frequently executed + short-running* operations
- * try grouping more user-managed tasks on the same WH → WH consolidation
- * *user-managed* (w/ WH) vs *serverless* tasks (no WH, 1.2x more expensive)
 - `USER_TASK_MANAGED_INITIAL_WAREHOUSE_SIZE` → **def MEDIUM!**
- * information_schema table functions
 - `SERVERLESS_TASK_HISTORY()` → for serverless tasks
 - `TASK_HISTORY()` → for user-managed tasks
 - `WAREHOUSE_METERING_HISTORY()` → for warehouses

*Replace Snowpipe with
Snowpipe Streaming*

Tip #64

CREDITS COST



Checkpoints

- * **Snowpipe + Snowpipe REST API**

- old legacy service + API for streaming in small batches

- * **Snowpipe Streaming API**

- more cost-effective data ingestion + save on storage costs
- high-throughput + low-latency streaming + low cost
- automated service (no VW)
- uses managed-compute resources + fixed credit charge per file

*Estimate Cost of
Automatic Clustering
on Tables*

Tip #65

Checkpoints

- * *Automatic Clustering* = serverless, in any edition
- * no clustering keys → no Automatic Clustering costs
- * Clustered Tables serverless price = 2 credits/hour
- * *AUTOMATIC_CLUSTERING_HISTORY* view/function
- * AUTOMATIC_CLUSTERING Snowflake WH
- * avoid table fragmentation → check load+query patterns

*Estimate Cost of the
Query Acceleration Service
(QAS)*

Tip #66

Cost

w/o QAS: 8min x 5c/min = **40c**

w/ QAS: 40sec x 5c/min x 14 nodes =< **47c**

X-Small WH (1 node) = \$3/h = 5c/min

QAS 14x scale factor =~ \$3/h/node = 5c/min/node

Programming QAS

- * `ENABLE_QUERY_ACCELERATION` - in WH, to enable
- * `QUERY_ACCELERATION_MAX_SCALE_FACTOR` - ctrls boost
- * `SYSTEM$ESTIMATE_QUERY_ACCELERATION` - w/ qID
- * `QUERY_ACCELERATION_ELIGIBLE` - system view
- * `QUERY_HISTORY` - system view, w/ data
- * `QUERY_ACCELERATION_HISTORY` - in info schema, for cost

Checkpoints

- * [Query Acceleration Service](#) (QAS) = serverless, on WH
- * can accelerate parts of a WH query workload
- * parts of query processing → shared compute resources
- * billed per second, no minimum
- * no cost if enabled + no eligible queries
- * in some cases could cost less than scaling up the WH

*Estimate Cost of the
Search Optimization Service
(SOS)*

Tip #67

Search Optimization Service

- * for some lookup/analytical queries/views
- * for queries w/ filters/joins, running for seconds+
- * not on the primary cluster key
- * at least one column must have 100K+ distinct values
- * for numbers, date, string, variant, geography, binary etc (not float, geometry)
- * filtering predicates: EQUALITY, SUBSTRING, GEO → keep a search access path
- * for IN, reg exp, substrings, JSON searches, geospatial queries, and/or
- * not for external/dynamic tables, MVs, casts...

Checkpoints

- * Enterprise+, 10 credits/h + 5 credits/h cloud (2+1 in sec. dbs)
- * avoid in high churn tables
- * avoid when data changes a lot in the table
- * enable only on specific columns, w/ the right predicate
- * disable once you no longer have benefiting queries
- * check actual query improvements
- * pair the service with smaller warehouses to lower the cost

*Reduce
Materialized Views
Maintenance Cost*

Tip #68

Checkpoints

- * MV cost = 10 credits/h + 5 credits/h cloud (2+1 in sec. dbs)
- * there are no tools to estimate costs of maintaining MVs (serverless)
- * *what may trigger MV maintenance*
 - any change of micro-partitions in the base table
 - any changes w/ DML statements on the base table
 - reclustering of the base table
- * cost proportional to: number of MVs + clustered MVs
- * suspend/resume MVs → only defers costs, not reducing them
- * *best practice*: start by creating only a few MVs + monitor the costs over time

*Reduce Database
Replication Cost*

Tip #69

Replication in Snowflake

- * **Database Replication** = ~db backup, for all editions
 - w/ replication groups, across multiple accounts
- * **Share Replication** = for all editions
- * **Account Replication** = Business Critical only, same org
 - w/ failover/failback ← w/ failover groups
- * **Cross-Cloud Auto-Fulfillment** = shared listing → other regions

Checkpoints

- * *db replication compute costs* = 2 credits/h managed + 1 credit/h cloud
- * standard costs in both source+target accounts!
- * initial+refresh replication → lower refresh frequency
- * avoid transfer between regions → data transfer costs
- * avoid very large databases, or w/ lots of data changes
- * refresh granular objects → use a replication group
- * DATABASE_REPLICATION_USAGE_HISTORY view → CREDITS_USED

*Estimate Cost
of Hybrid Tables*

Tip #70

Hybrid Tables

- * **Snowflake Unistore** = OLTP (transactional) + OLAP (analytical)
- * fast writes as well → stores data also by row
- * required primary keys
- * optional foreign keys → referential integrity enforced
- * allows indexes
- * row locking

Checkpoints

- * Standard+, PuPr in some AWS paid accounts (currently)
- * *consumption modes*
 - **compute** = for WHs, same credits as for standard tables
 - **storage** = fix monthly rate per GB, more expensive
 - **requests** = for serverless resources on the row storage clusters
- * *cost per request*
 - 1 credit/h managed + 1 credit/h cloud
 - 1 credit per 30 GB read + 1 credit per 7.5 GB write

Introduction to Data Storage

Tip #71: Use On-Demand When You Don't Know Your Spending Pattern

Tip #72: Copy and Keep Less Data

Tip #73: Lower Data Retention with No Time Travel

Tip #74: Estimate Storage Cost of the Fail-Safe

Tip #75: Use Transient or Temporary Tables

Tip #76: Use Zero-Copy Cloning

Tip #77: Clone Less Data

Tip #78: Ensure Tables Are Clustered Correctly

Tip #79: Drop Unused Tables and Other Objects

Tip #80: Remove Old Files from Stage Areas

*Use On-Demand Storage
When You Don't Know
Your Spending Pattern*

Tip #71

Data Storage Costs

OPTIMIZED STORAGE

Snowflake charges a monthly fee for data stored in the platform. Calculated using the average amount of storage used per month, after compression, for data ingested into Snowflake.

[Download Snowflake Pricing Guide](#) 



On-Demand Storage

Pay for usage month-to-month.

\$23 / per TB / per month (\$USD)

AWS, US West (Oregon)



Capacity Storage

Discounted storage paid upfront.

\$23 / per TB / per month (\$USD)

AWS, US West (Oregon)

Capacity Storage Pricing

Table 3(a): Standard Storage Pricing

Cloud Provider	Region	On Demand Storage Pricing (TB/mo)	Capacity Storage Pricing by ACV Range (TB/mo)						
			Tier 1	Tier 2	Tier 3	Tier 4	Tier 5	Tier 6	Tier 7
			USD \$0 - \$1,199,999	USD \$1,200,000 - \$2,999,999	USD \$3,000,000 - \$4,999,999	USD \$5,000,000 - \$9,999,999	USD \$10,000,000 - \$19,999,999	USD \$20,000,000 - \$39,999,999	USD \$40,000,000+
AWS	US East (Northern Virginia)	\$23.00	\$23.00	\$21.47	\$19.94	\$18.40	\$16.86	\$15.34	\$13.80
AWS	US West (Oregon)	\$23.00	\$23.00	\$21.47	\$19.94	\$18.40	\$16.86	\$15.34	\$13.80

Checkpoints

- * On-Demand (PAYG) vs Capacity (prepaid) Storage Pricing
 - US \$23/TB/mo → lowest for AWS/Azure US (\$20 for GCP Central US)
 - storage billing → average stored per day
- * Capacity Storage → less only for \$2M+ accounts
 - storage price → per TB/mo (from highest for On-Demand)
 - long-term contract for small discounts → could overspend
 - auto-switched to On-Demand if underspent
 - use eventually once you know your spending patterns

*Copy and Keep
Less Data*

Tip #72

Cost

- * table data storage: 2B rows (65GB)
- * Time Travel: 1 day (def)
- * Fail-safe: 7 days (fixed)
- * one clone: 2B rows (shared)

copied from tpch_sf100tbl.store_sales (288B rows, 10TB) in 11 mins

TABLE_STORAGE_METRICS

- * in INFORMATION_SCHEMA → w/ 1-2h latency
- * **Active bytes** = live data that can be queried
- * **Bytes in Time Travel** = deleted, but still kept for the retention period
- * **Bytes in Fail-safe** = deleted, past the Time Travel, but still kept for 7 days
- * **Bytes retained for clones** = deleted, but still kept as clone refs

Retention Policies

- * do not use Snowflake as an archiving system
 - backup/store raw data cheap in Amazon S3 Glacier
- * ex: retention policy: 5 → 3 years
 - less bytes scanned in queries
 - 36% storage cost saving

Checkpoints

- * copy less, clone more → for dev/test/prod envs
- * use transient/temporary tables for no Fail-safe
- * disable Time Travel
- * cloned tables will transparently keep shared table data
- * create huge table + delete right away → still expensive!
- * limit period of retention policy

*Lower Data Retention
with No Time Travel*

Tip #73

Checkpoints

- * *permanent table*

- 0..1 days (def 1) Time Travel for Standard Edition
- 0..90 days (def 1) Time Travel for Enterprise Edition

- * *temporary/transient table*


- 0..1 days (def 1) Time Travel

Time Travel may remain there after you drop a table
set `DATA_RETENTION_TIME_IN_DAYS = 0` to disable

*Estimate Storage Cost
of the Fail-Safe*

Tip #74

Fail-Safe for Large Dropped Table



Cost

- * 1TB table dropped after 1 day $\rightarrow \$23 / 30 \text{ days} \approx 77c$
- * Fail safe for dropped table $\rightarrow \$23 \times 7/30 \text{ days} = \5.37

cost On-Demand storage: US \$23/TB/mo


Checkpoints

- * permanent table : 7 days Fail-safe
- * temporary/transient table: no Fail safe
- * drop a large table → *all* its data moves into Fail safe!

*Use Transient
or Temporary Tables*

Tip #75

ELT Data Pipeline




Checkpoints

- * **temporary table** = auto-deleted at the end of session
- * **transient table** = shared, req explicit drop
- * max time travel 1 day (def 1) → save on storage
- * no fail safe (else 7 days) → save on storage
- * for transitory data (no need for backup)

*Use
Zero-Copy
Cloning*


Tip #76

Table Copy



Storage Cost for Copied Table = Original Data + New Data - Deleted Data

Table Clone



Storage Cost for Cloned Table = Changed Data + New Data

Checkpoints

- * **copied table cost** = orig data + INSERTs - DELETES
- * **cloned table cost** = UPDATES + INSERTs
- * clones will save on storage (keep data references)
- * clones will save on compute (to duplicate data)
- * clones will avoid redundancies → less error prone

*Clone
Less Data*

Tip #77

Checkpoints

- * extensive DDL (cloning) → significant cloud serv usage
- * DDL ops (like cloning) use only cloud compute serv
- * clone more granular objects
- * clone only individual tables rather than entire schemas

*Ensure Tables Are
Clustered Correctly*

Tip #78

CLUSTERING_INFORMATION

- * *cluster_by_keys*, *total_partition_count* - just info
- * *total_constant_partition_count* - higher is better (reclustered)
- * *average_overlaps* - lower is better (<1), higher not well-clustered
- * *average_depth* - lower is better (CLUSTERING_DEPTH)
- * *partition_depth_histogram*
 - 0, 1, 2, 3, ..., 16, 32, 64, 128, ...
 - most w/ depth 1 or on top is best
 - worst if most at the bottom
- * *clustering_errors* - max 10 recent reclustering errors [...]

Checkpoints

- * load data → natural auto-clustering
- * evaluate for natural auto-clustering (no key → no cost!)
- * evaluate for a clustering key (auto-clustering cost!)
- * identify + evaluate for candidate columns
- * load data according to current/new clustering key
- * use queries that benefit from current table clustering

*Drop Unused Tables
and Other Objects*

Tip #79

Storage Usage History

- * **DATABASE_STORAGE_USAGE_HISTORY**

- *average_database_bytes*
- *average_failsafe_bytes*
- *average_hybrid_table_storage_bytes*

- * **STAGE_STORAGE_USAGE_HISTORY**

- *average_stage_bytes*

Table Storage Metrics

* *active_bytes*

* *time_travel_bytes*

* *failsafe_bytes*

* *retained_for_clone_bytes*


Checkpoints

- * *storage*
 - db active + time travel + fail-safe + retained for clones + hybrid tables + stages
- * drop any unused objects → recreate them from scripts!
- * *unused tables* = not modified/changed/accessed lately
- * **TABLES** view
 - *last_altered* → by DDL/DML operation
- * **ACCOUNT_HISTORY** view
 - *base_objects_accessed* → JSON array, must flatten!
 - *query_start_time* → last accessed time

*Remove Old Files
from Stage Areas*

Tip #80

Storage Wasted by Old Stage Files



Checkpoints

- * delete files no longer required from stages
- * external+internal stages (named/table/user)
- * ACCOUNT_USAGE.STORAGE_USAGE.**stage_bytes**

SHOW STAGES IN ACCOUNT → all stage names

LIST @mystage → all files

REMOVE @mystage / @%mytable / @~

Introduction to Data Transfer

Tip #81: Data In is Free, Data Out is Expensive

Tip #82: Choose the Same Provider and Region Where Your Data Is

Tip #83: External Access Integrations vs External Functions

Tip #84: Use Data Compression

Tip #85: Use Batch Transfer with Path Partitioning

Tip #86: Use Bulk Loads instead of Single-Row Inserts

Tip #87: Use Parallel Data Uploading

Tip #88: Design Cost-Effective Data Pipelines

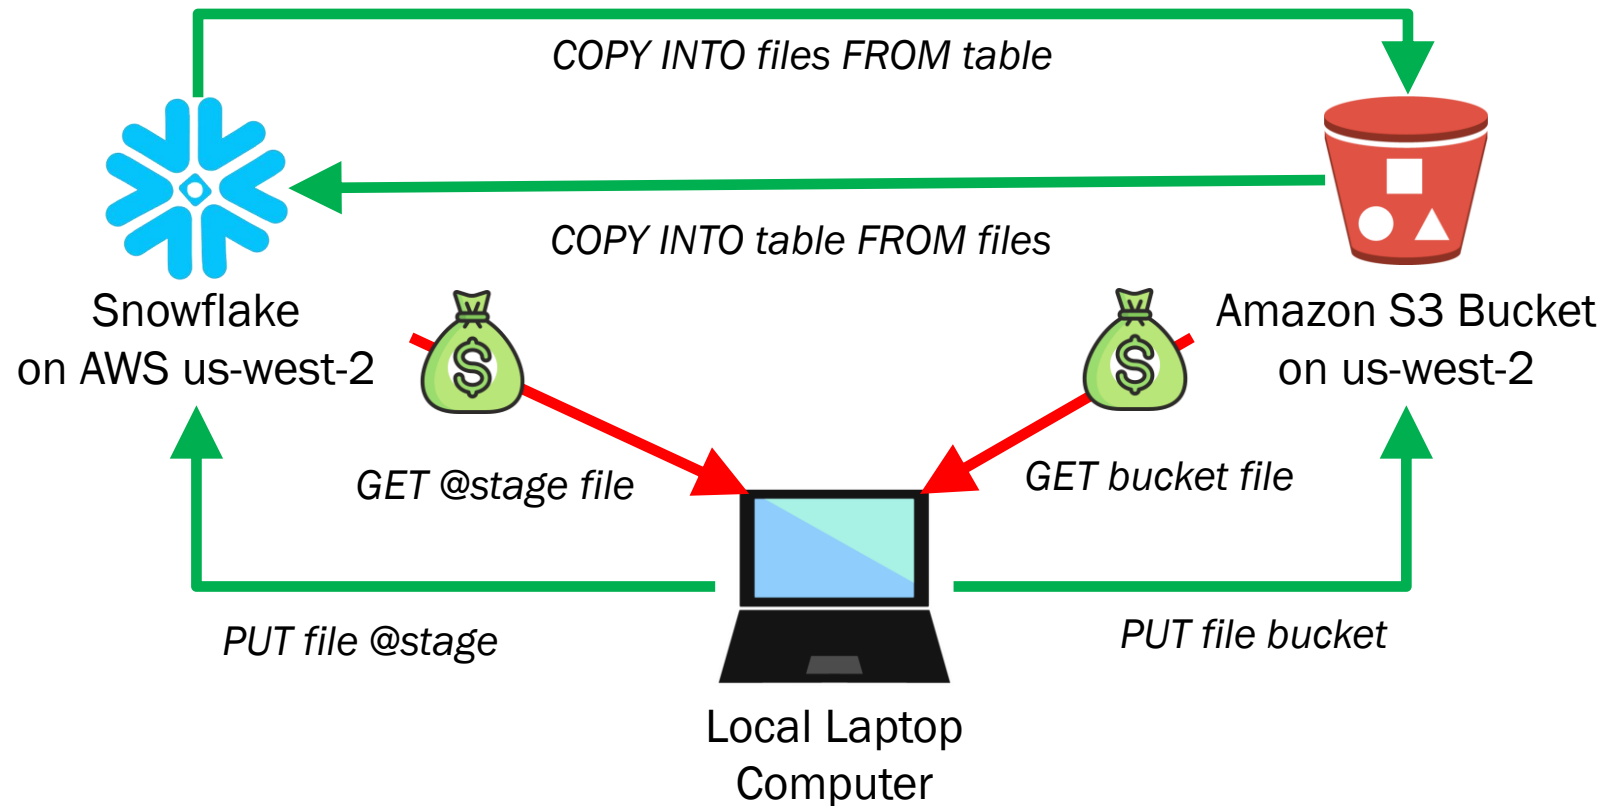
Tip #89: Use External Tables in a Data Lake

Tip #90: Query Parquet Files instead of CSV

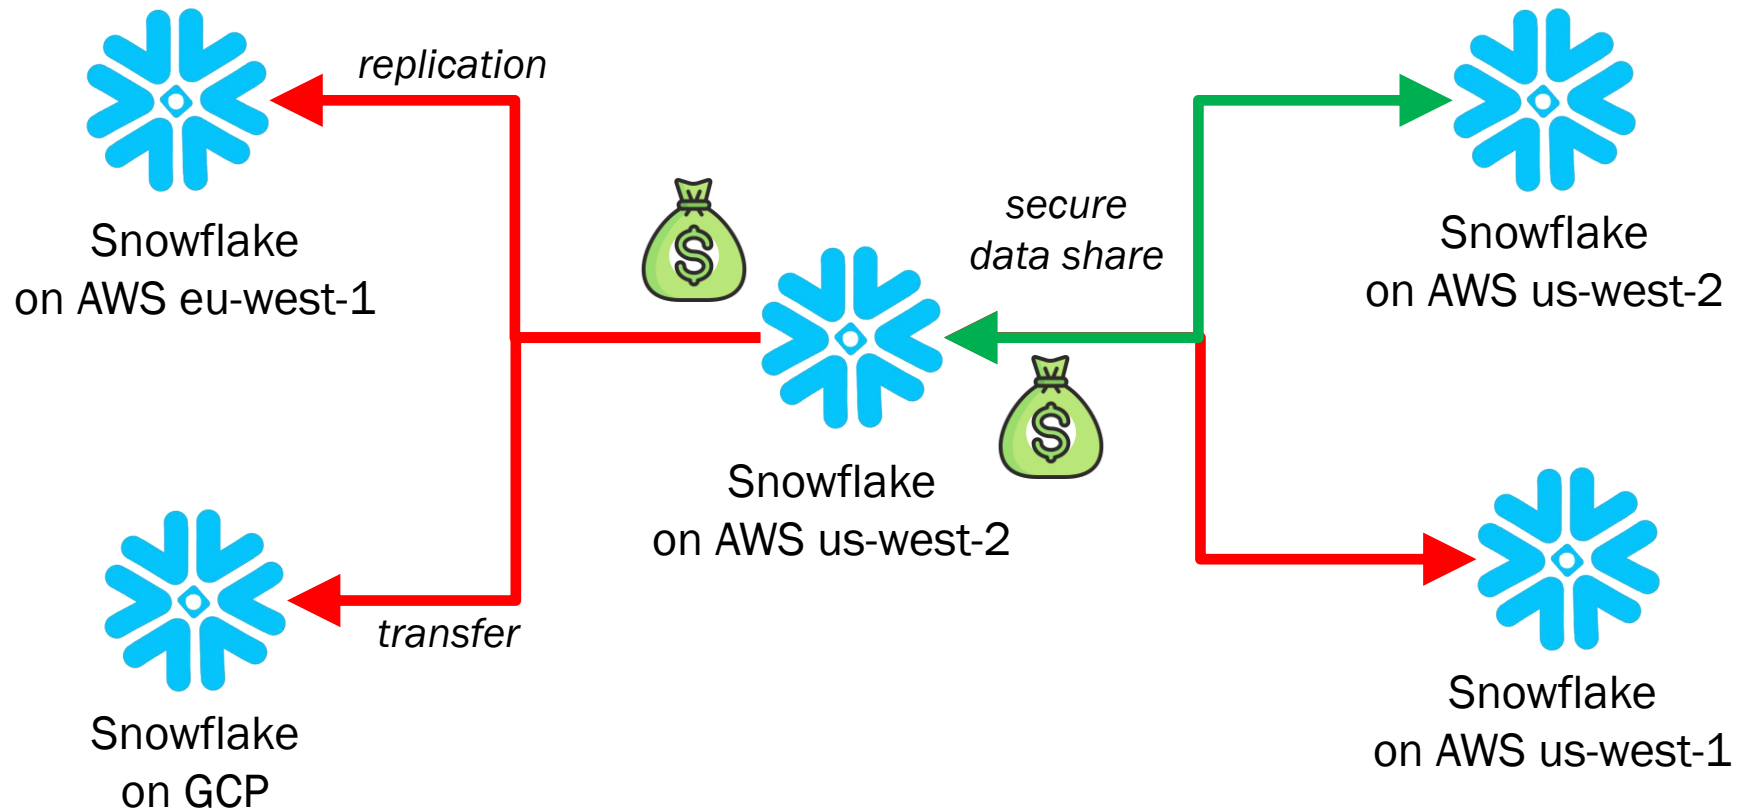
*Data In is Free,
Data Out is Expensive*

Tip #81

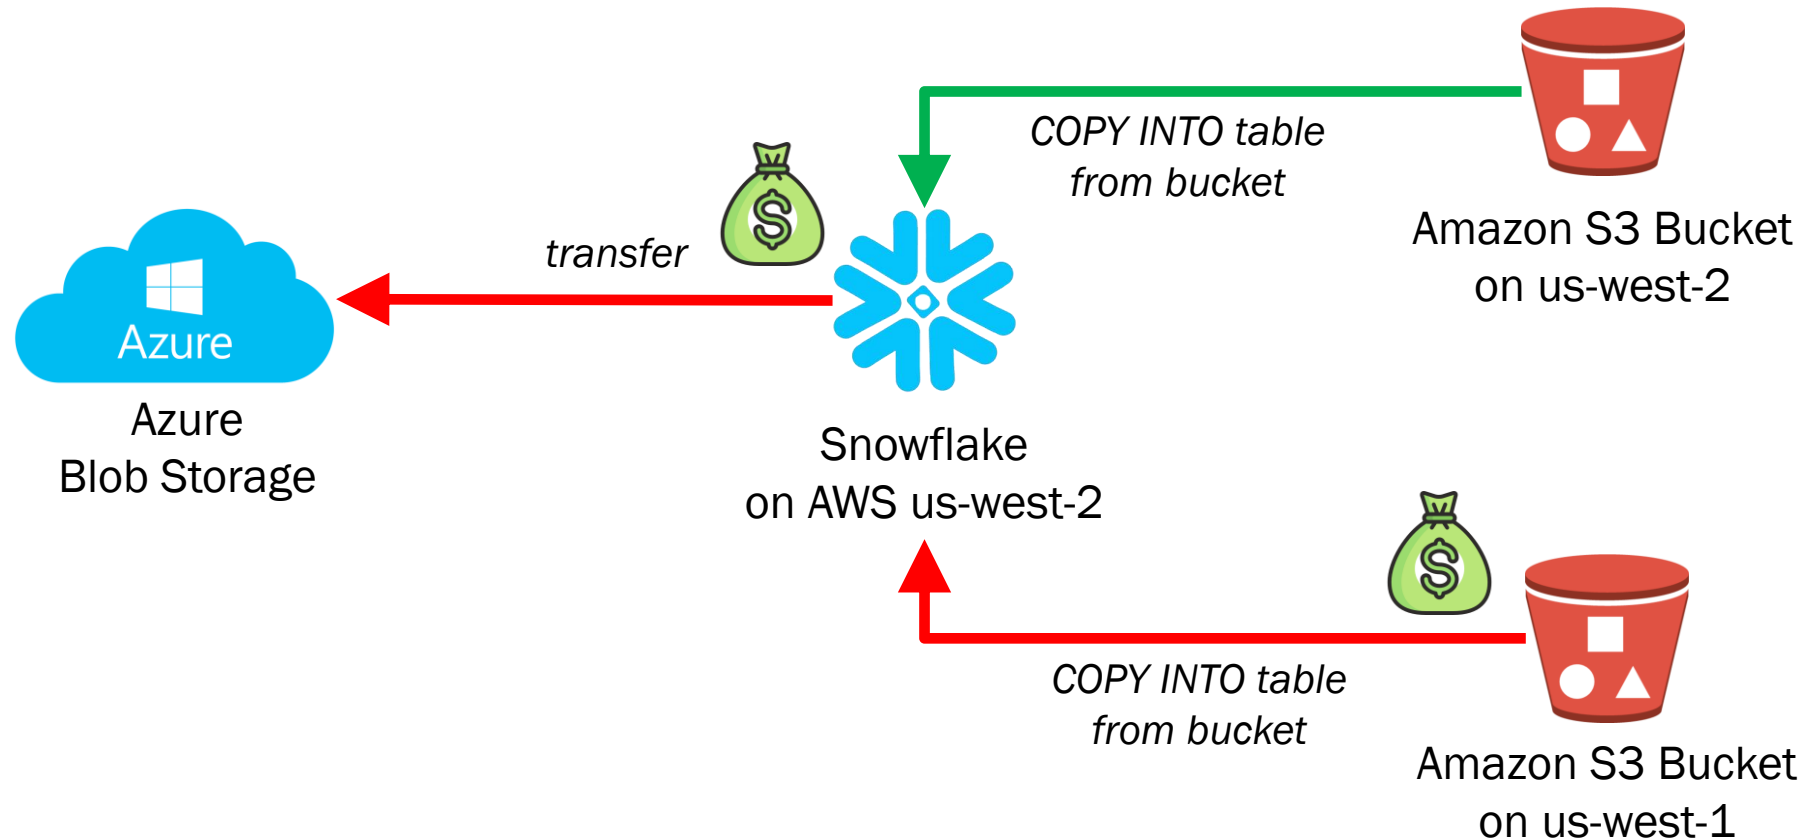
Local Data Transfer Costs



Snowflake-to-Snowflake Costs



Internet Data Transfer Costs



DATA_TRANSFER_HISTORY

	TFX ...	TRANSFER_TYPE	BYTES_TRANSFERRED
1	aws (us-west-2) → (aws) us-west-2	EXTERNAL_FUNCTION	2808091
2	aws (us-west-2) → (aws) us-east-2	EXTERNAL_FUNCTION	27797
3	aws (us-west-2) → (internet) internet	EXTERNAL_ACCESS	98058
4	aws (us-west-2) → (aws) us-west-2	COPY	101584

Checkpoints

- * Snowflake (us-west-1) → S3 bucket (us-west-2) - \$ to Snowflake
- * Snowflake (us-west-1) ← S3 bucket (us-west-2) - \$ to AWS
- * Snowflake → local - \$ to Snowflake
- * AWS → local - \$ to AWS

*Choose the Same
Provider and Region
Where Your Data Is*

Tip #82

Data Transfer Pricing

- * **AWS**: diff region \$20+ / diff prov \$90+
- * **Azure**: same cont \$20+ / diff cont \$50+ / diff prov \$87.50+
- * **GCP**: same cont \$10+ / diff cont \$80+ / diff prov \$120+

pricing is per TB of data transferred

within the same region+provider is FREE

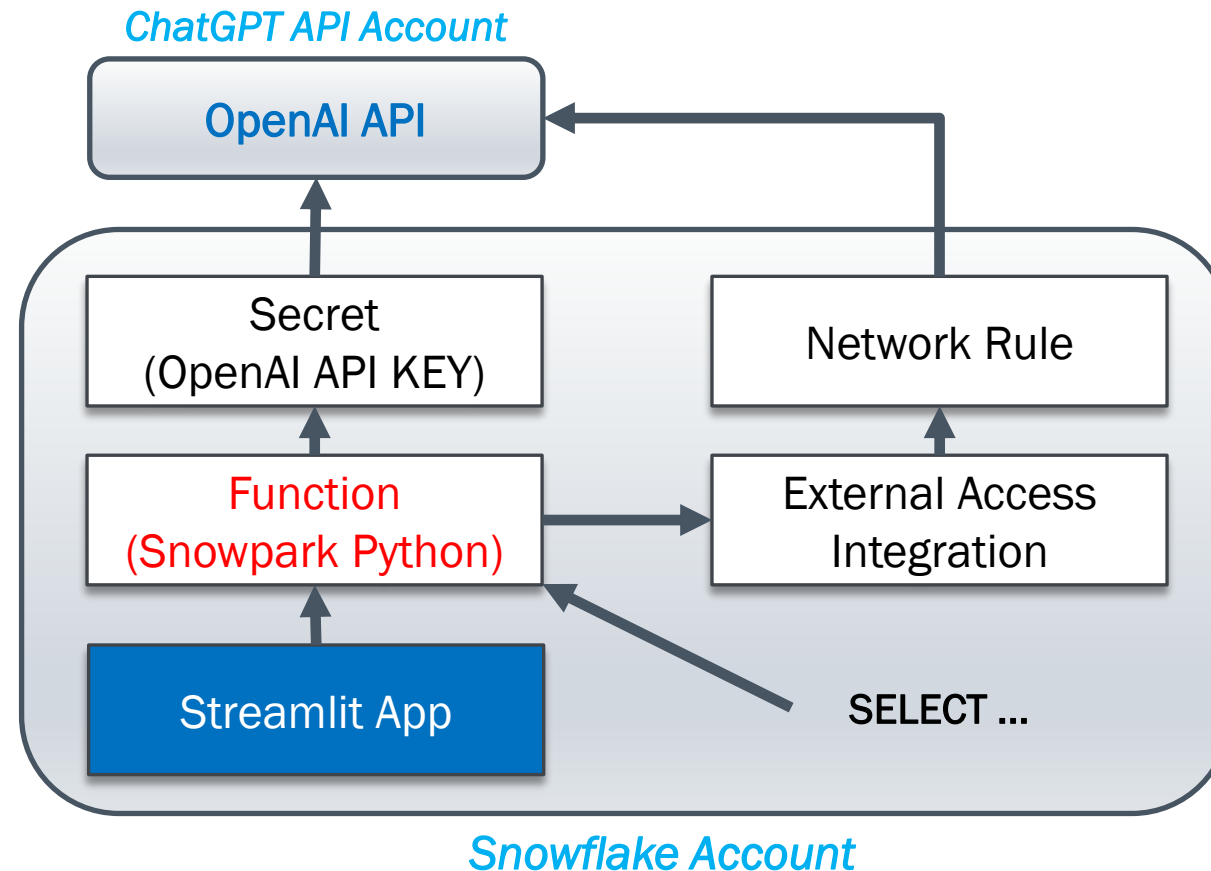
Checkpoints

- * no cost to tfx data in (upload)
- * no cost to tfx data out within the same region+provider
- * moderate cost to tfx to another region, same provider
- * moderate cost to tfx within the same continent+provider
- * higher cost to tfx to another continent, same provider
- * highest cost to tfx to another provider

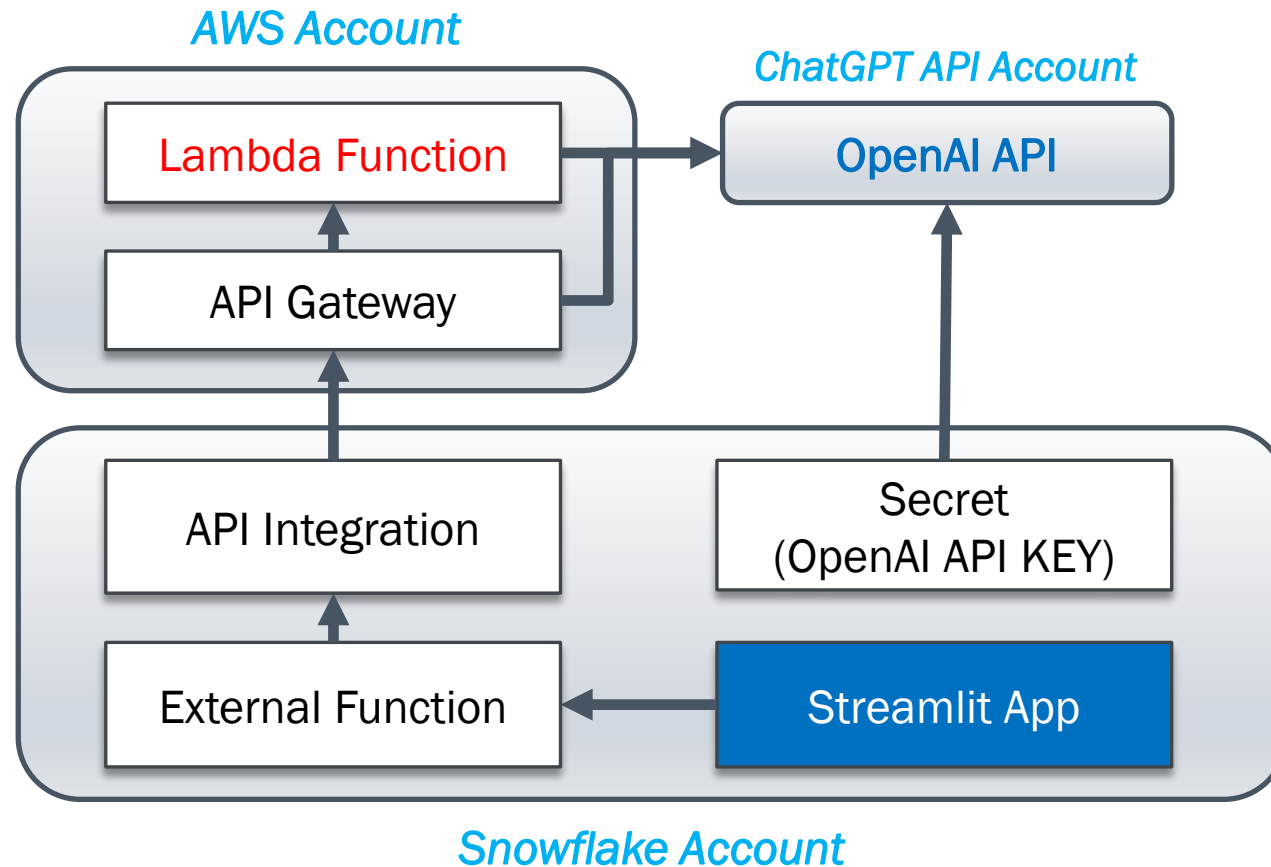
*External Access
Integrations
vs External Functions*

Tip #83

External Access Integration



External Function



Checkpoints

* *External Access Integration*

- business logic mostly within Snowflake, in the WH
- use when you do not transfer a lot of data w/ the remote API
- easier to implement and deploy

* *External Function*

- business logic mostly remote, in another cloud service
- use when you may transfer a lot of data w/ the remote API
- more complex to implement and deploy


Use

Data Compression

Tip #84

LOADING STRUCTURED DATA

■ load time (seconds)



Checkpoints

* Snowflake storage

- min 3:1 compression ratio → billed only compressed storage!
- 4.7 TB uncompressed (28.8B rows) → stored as 1.3 TB

* loading compressed data

- less data transfer cost + less time (cheaper COPY INTO)
- from CSV (GZIP) >> faster than Parquet/ORC (Snappy)
- keep data files 100-250 MB in size compressed
- COMPRESSION = GZIP/BZ2/BROTLI/ZSTD/DEFLATE/RAW_DEFLATE

*Use Batch Transfer
with Path Partitioning*

Tip #85

Checkpoints

- * partition staged data files w/ logical granular paths
- * COPY INTO must list S3 folder files → improve performance!
- * ex: <s3://bucket/app2/loc1/2016/07/01/14/>
- * copy any fraction of data w/ a single command
- * could copy data by hour/day/month etc
- * create subfolders of 10-15 mins increments per hour

*Use Bulk Loads
instead of
Single-Row Inserts*

Tip #86


Checkpoints

- * **COPY INTO** (batch/bulk) vs **INSERT** (single row)
- * OLTP vs OLAP
- * batch/bulk copy consumes less cloud res + faster
- * batch/bulk loads can be done in parallel
- * not for hybrid tables (Snowflake Unistore)

*Use Parallel
Data Uploading*

Tip #87

Parallel Data Uploading



Checkpoints

- * small WH + one file → ok in most cases
- * big WH + one file → may take longer if large file
- * big WH + single file → could waste idle nodes
- * big WH + parallel loading → faster for ~100 files

```
COPY INTO mytable FROM @stage/folder/ FILES=('file1.csv', ...)
```

*Design Cost-Effective
Data Pipelines*

Tip #88

Checkpoints

- * check that data-loading processes don't run more often than consuming data pipelines
- * reduce sync frequency of your data ingestion tool + evaluate how often to update data
- * upload data to S3 + use as external stage for data loading to Snowflake
- * check related costs when running all type of workloads in Snowflake
- * it could be more cost-effective to transform S3 data w/ AWS Glue, then load it in Snowflake
- * incremental processing: do not rerun entire batches of data for minor updates
- * restart interrupted pipelines from the point of interruption, w/ no redundant computation

Use
External Tables
in a Data Lake

Tip #89

Checkpoints

- * to query in-place data lake files, from ext stages, R/O
- * slower than Snowflake tables → use mat view
- * metadata\$filename + metadata\$file_row_number
- * value → single VARIANT column, use for column defs (or INFER_SCHEMA)
- * partitioning always recommended → PARTITION BY col1, ...
- * manual/auto-refresh metadata → w/ event notifs, may incur cost
- * PIPE_USAGE_HISTORY → estimate charge

*Query Parquet Files
instead of CSV*

Tip #90

Parquet vs CSV

- * 80%+ less data stored
- * 30%+ faster queries
- * 99% less data scanned
- * 99%+ cost savings

Dataset	Size on Amazon S3	Query Run Time	Data Scanned	Cost
Data stored as CSV files	1 TB	236 seconds	1.15 TB	\$5.75
Data stored in Apache Parquet Format	130 GB	6.78 seconds	2.51 GB	\$0.01
Savings	87% less when using Parquet	34x faster	99% less data scanned	99.7% savings

Checkpoints

- * transform CSV to Parquet when left in the data lake
- * less storage because of binary format + compression
- * less data scanned because of less storage + columnar
- * faster queries because of the columnar format
- * huge compute + storage cost savings

Introduction to Snowflake Apps

Tip #91: Estimate Cost Impact of Data Sharing in Snowflake

Tip #92: Estimate Cost Impact of Client and Server (Snowpark) Apps

Tip #93: Estimate Cost Impact of Streamlit in Snowflake and Native Apps

Tip #94: Estimate Cost Impact of Data Science Applications

Tip #95: Check All Connected Applications

Tip #96: Third-Party Apps Saving Money Will Spend Money

Tip #97: Free Marketplace Native Apps Will Cost Money

Tip #98: Keep App Versions Updated

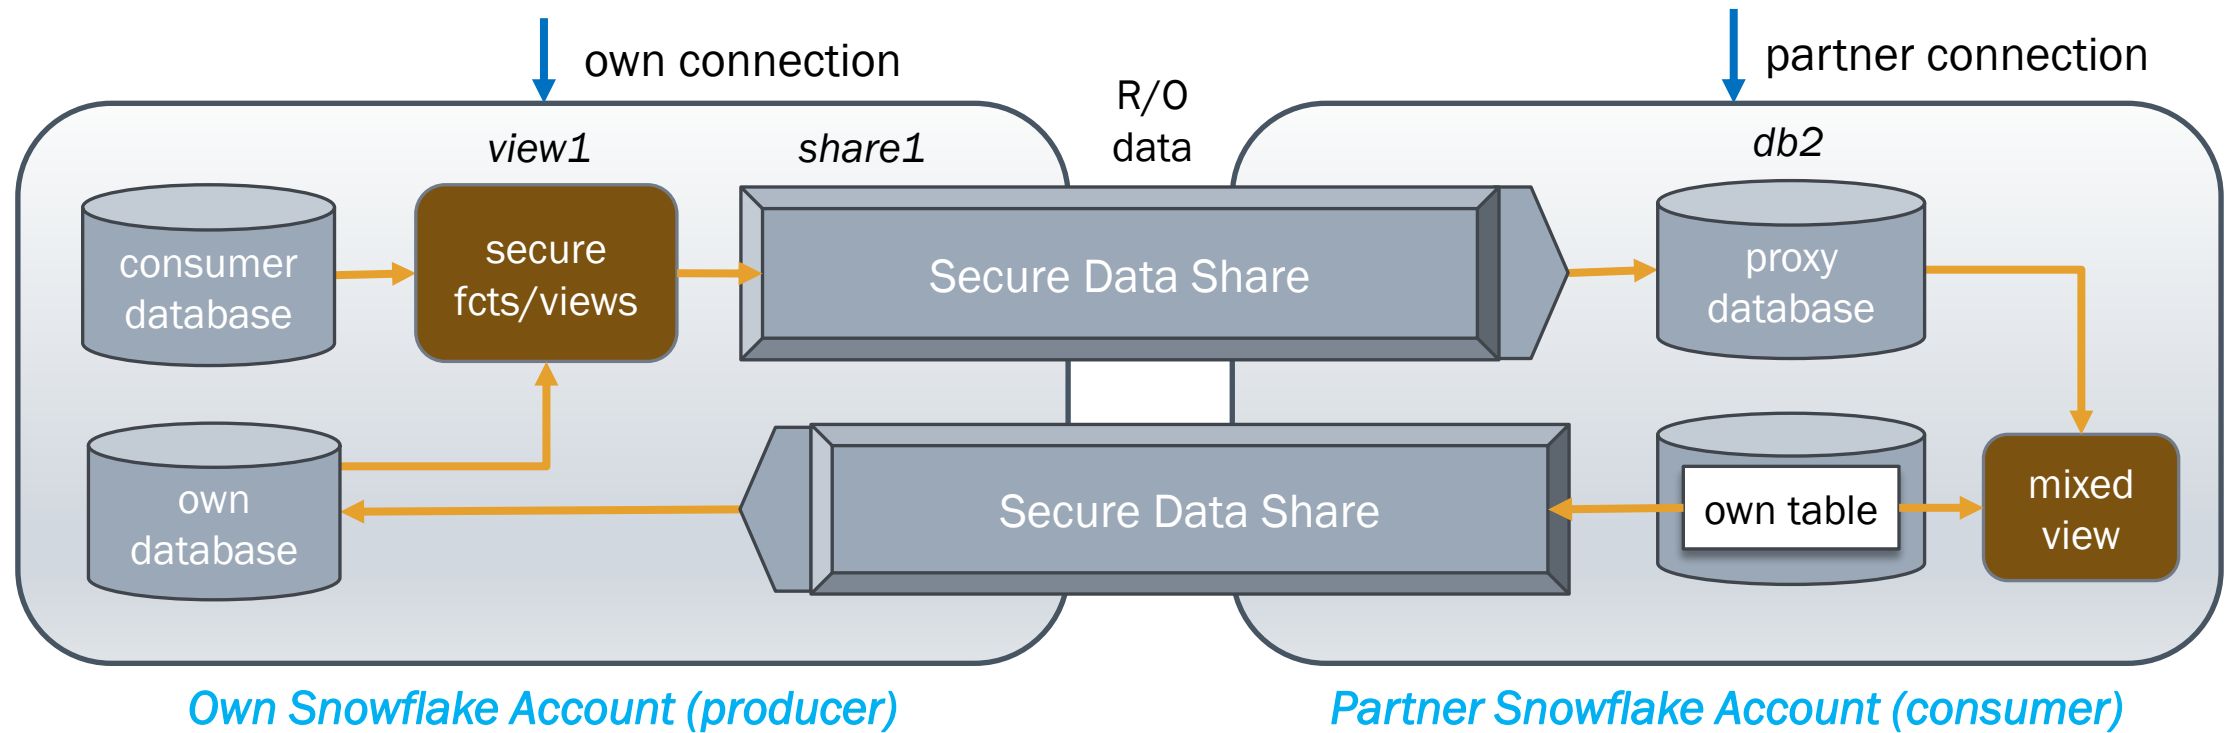
Tip #99: Cache Data in Third-Party Tools

Tip #100: Auto-Abort Running Queries from Disconnected Apps

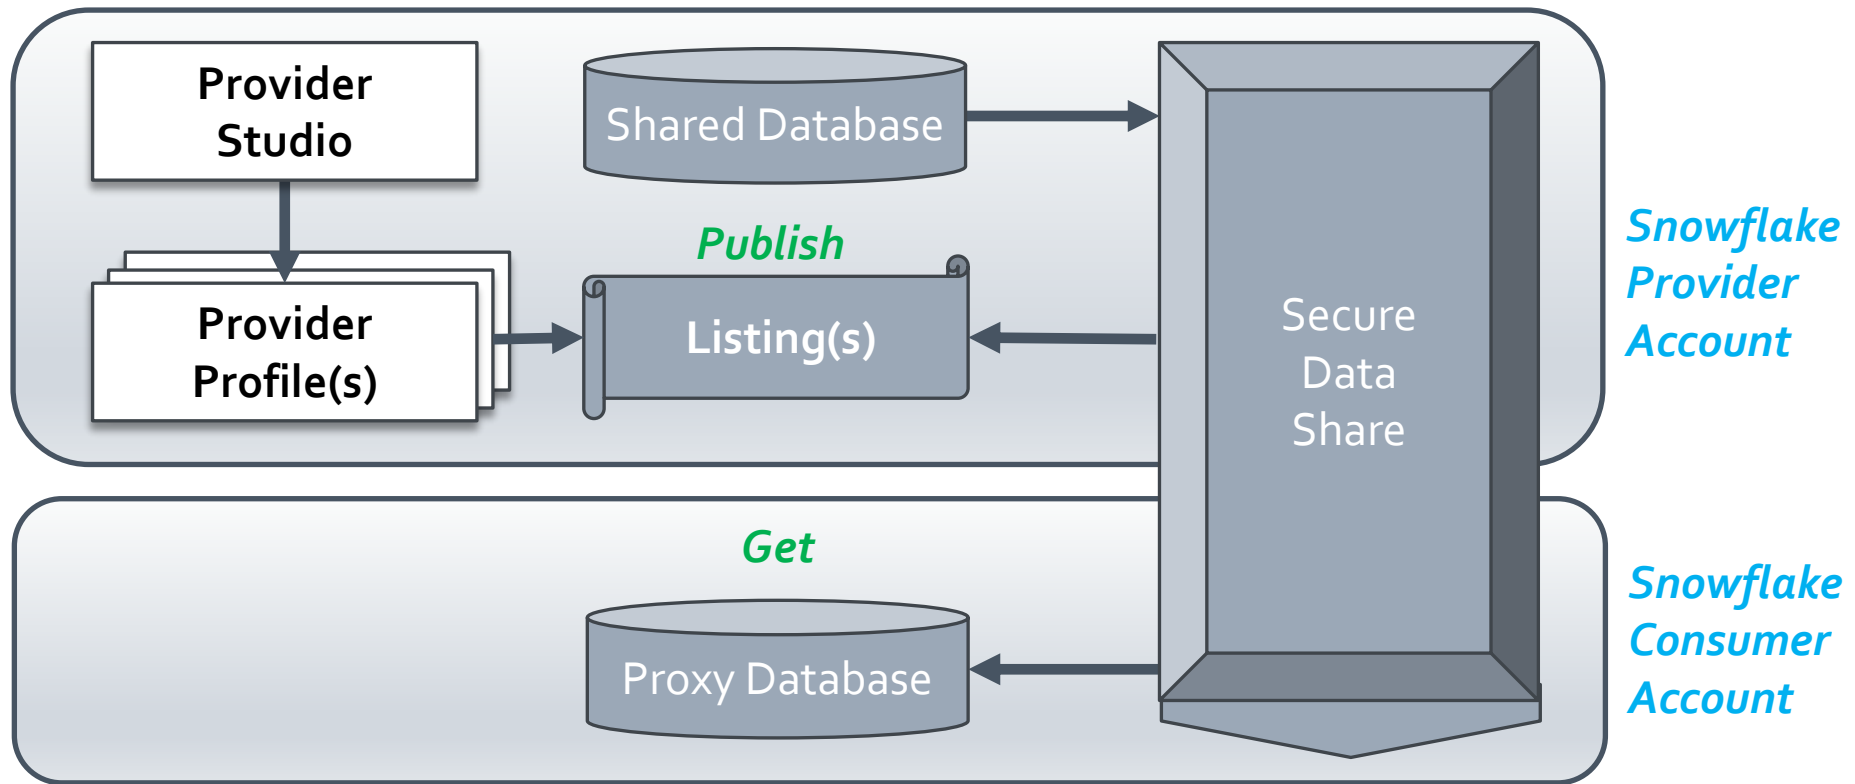
*Estimate Cost Impact
of Data Sharing
in Snowflake*

Tip #91


Inbound/Outbound Data Shares



Private Data Exchange



Data Share w/ Reader Account



Checkpoints

* data shares

- if provider → consume some other account compute credits
- if provider for Reader account → consume OWN credits!
- if consumer → consume OWN credits


* Marketplace (public) / Data Exchange (private)

- consumer of shared datasets → consume OWN compute credits
- ex: SNOWFLAKE_SAMPLE_DATA datasets → some are huge!

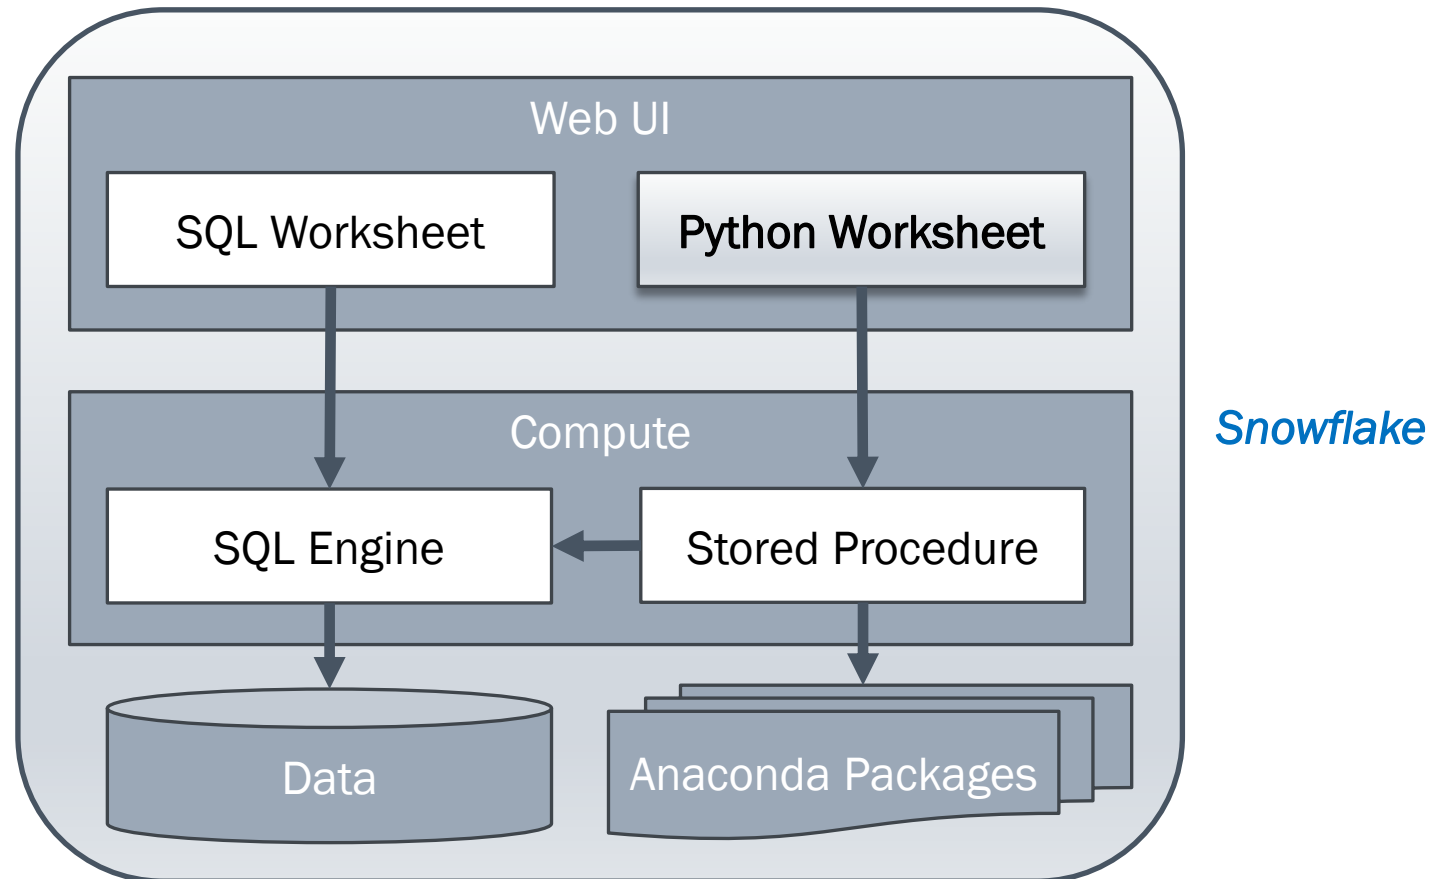
*Estimate Cost Impact
of Client and Server
(Snowpark) Applications*

Tip #92

Snowpark vs Python Connector



Python Worksheet



Client and Server Applications

- * **client apps** (w/ client drivers)
 - consume YOUR compute + storage credits
 - plus lots of data transfer, eventually
- * **server apps** (w/ Snowpark, SiS, Container Services)
 - consume **A LOT OF** YOUR compute + storage credits
 - but less data transfer (as they are closer to data)

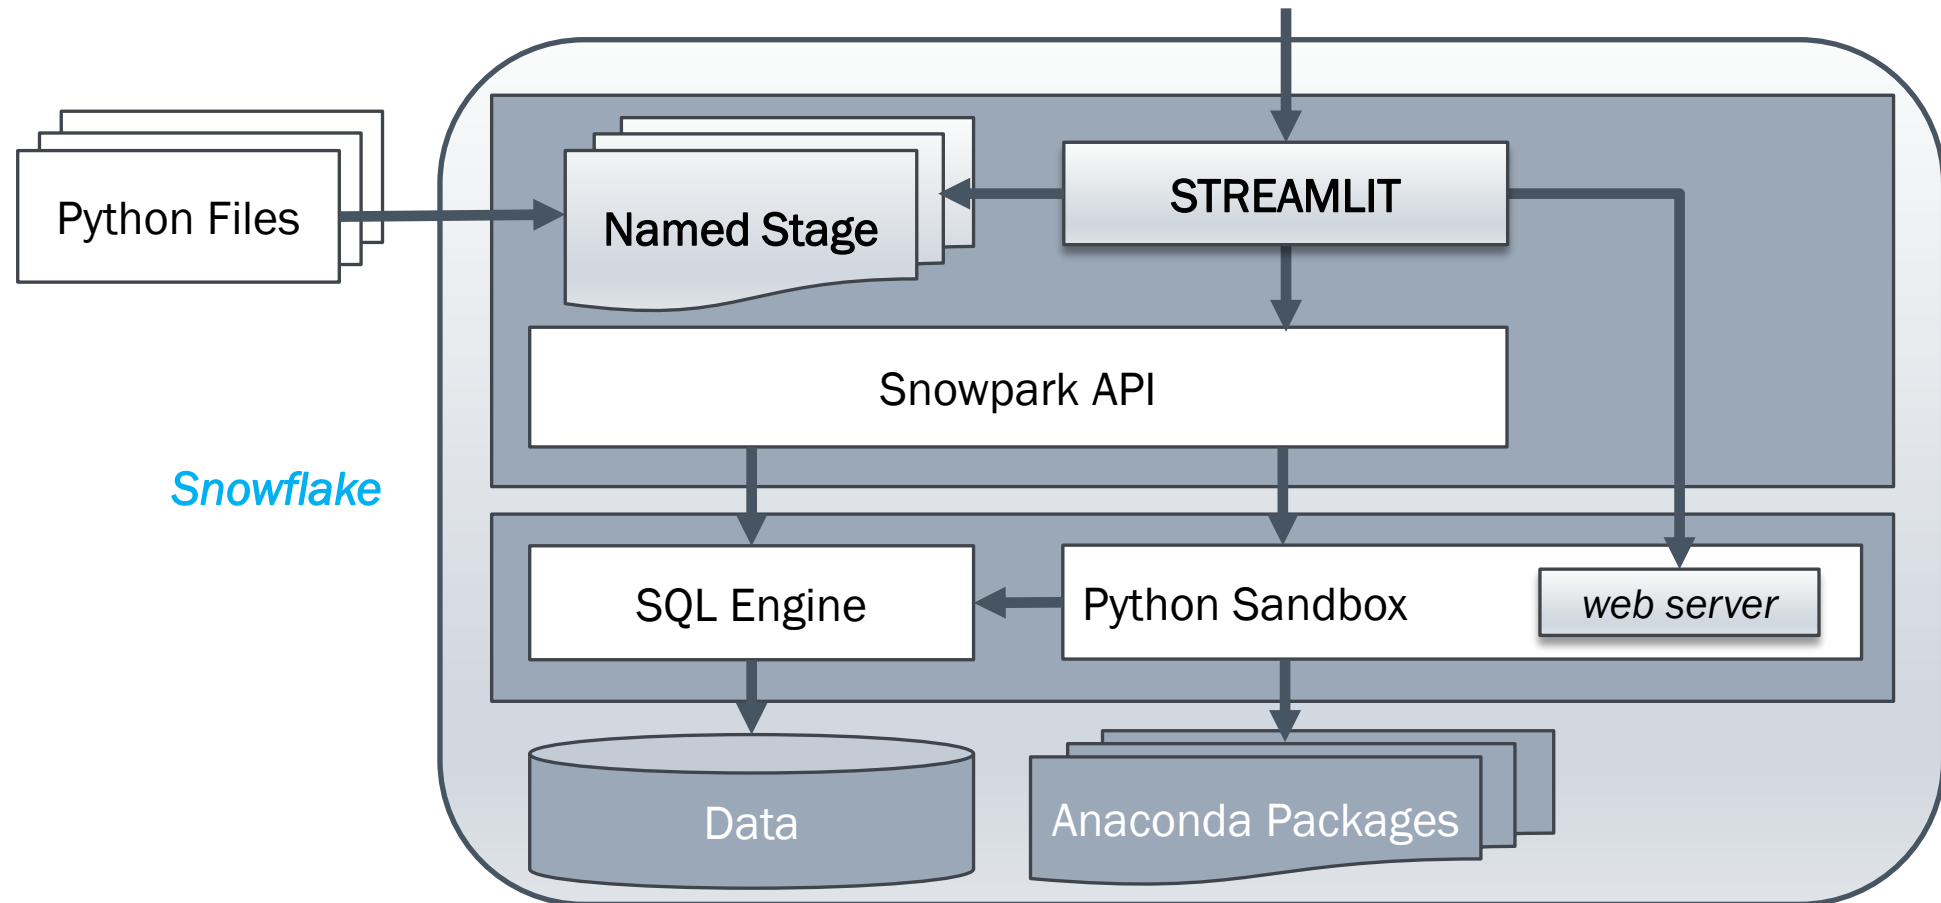
Checkpoints

- * Snowpark apps = mostly server-side apps
- * unlike client apps, business logic is deployed in WHs
- * Snowflake WHs = db servers (SQL) → app servers (Python)
- * better rich alternative to Snowflake Scripting
- * stored proc wrappers around Python code → in WHs
- * efficient if executing code closer to data is justified

*Estimate Cost Impact
of Streamlit in Snowflake
and Native Applications*

Tip #93

Streamlit in Snowflake (SiS)



Snowflake Native Apps

- * if provider → consume some other account credits
- * if consumer → consume OWN compute credits
- * *warning*: any FREE native app will cost you money!
- * ex: SNOWFLAKE app, w/ ACCOUNT_USAGE schema
- * it is your metadata, but expensive to browse!

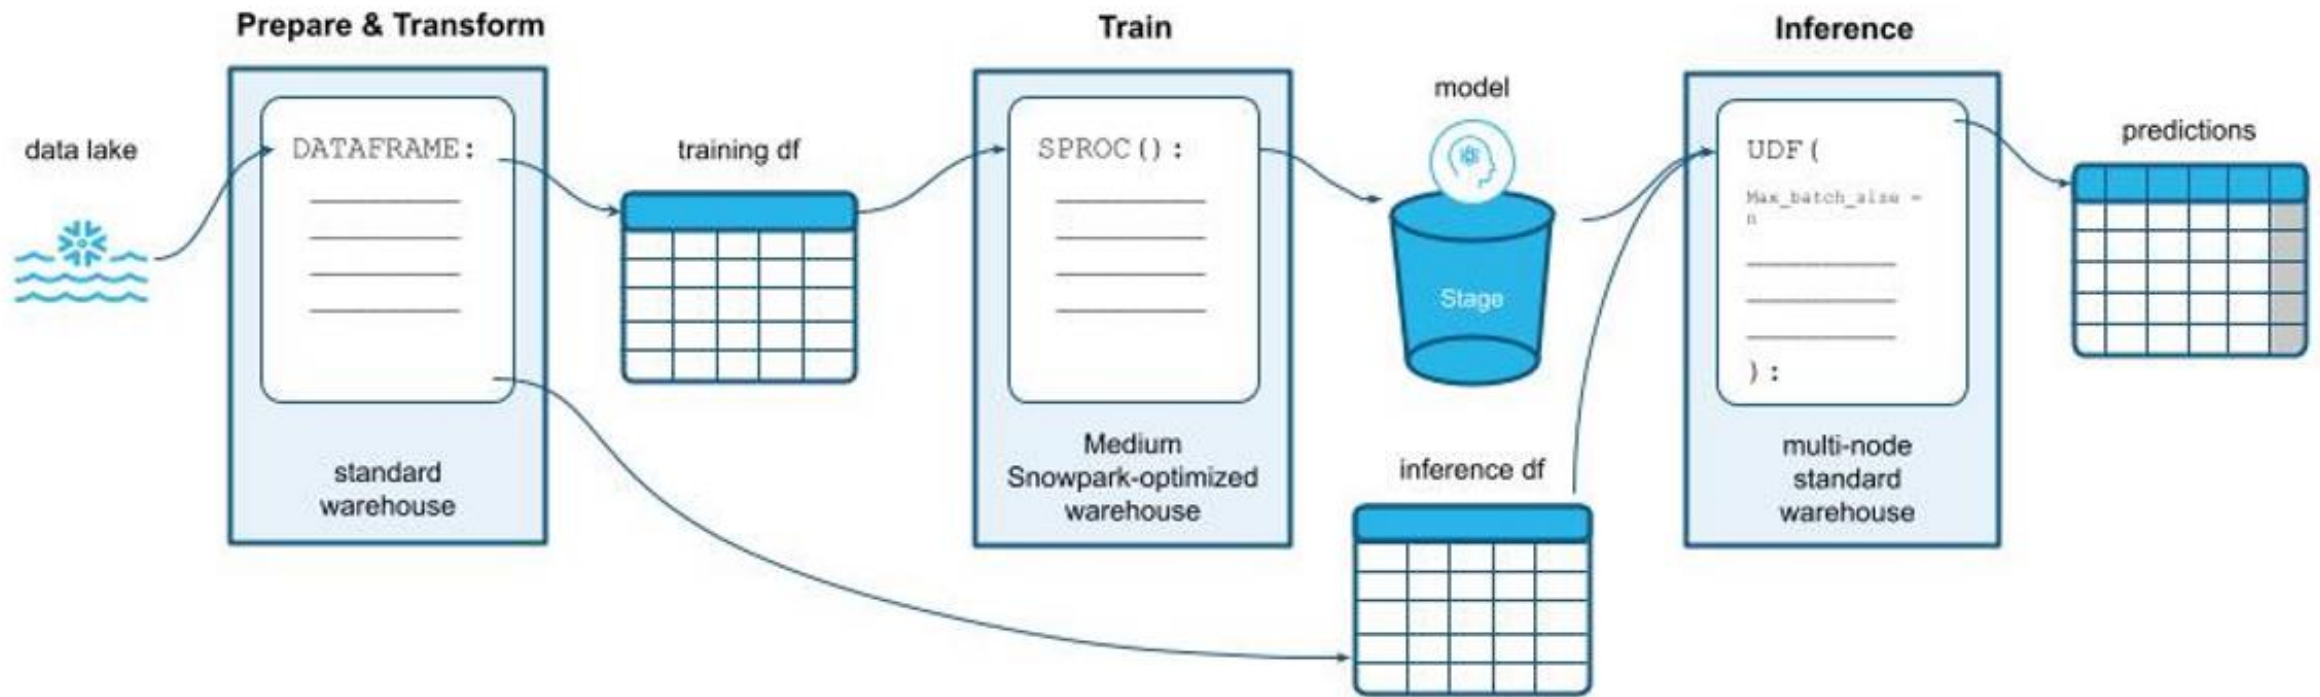
Checkpoints

- * **Streamlit in Snowflake (SiS)** = wrapper around Python code
- * **STREAMLIT** object, deployed+executed as stored proc in WH
- * it keeps WH up as long as app used → expensive!
- * app goes to sleep when WH auto-suspended
- * most Streamlit web apps can be easily deployed as SiS
- * alternatives: native apps, Container Services

*Estimate Cost Impact
of Data Science
Applications*

Tip #94

Model Training & Serving



<https://www.snowflake.com/blog/snowpark-python-feature-engineering-machine-learning/>

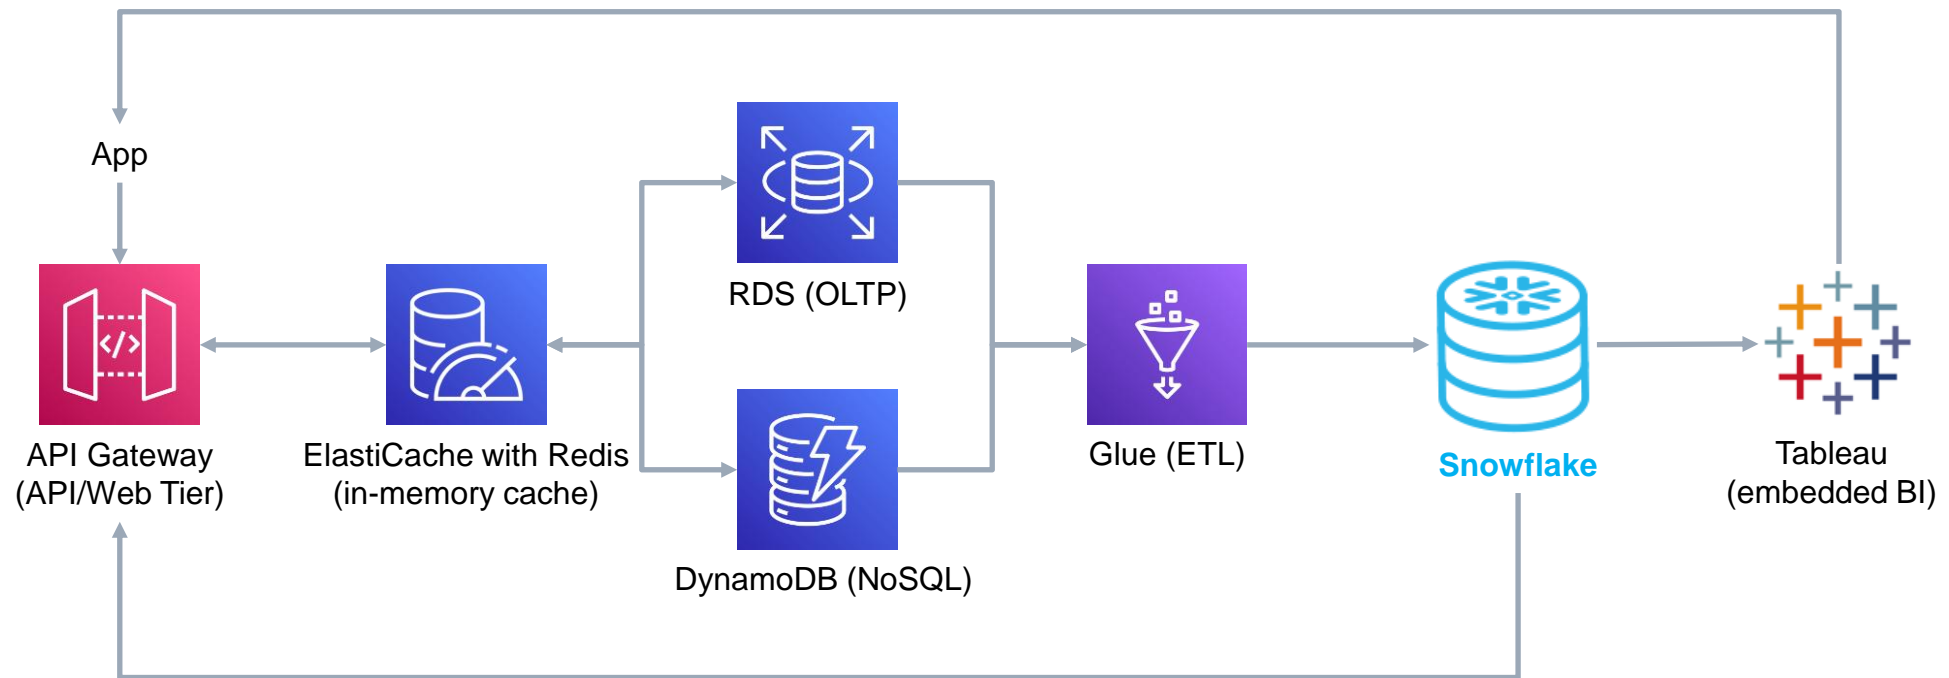
Checkpoints

- * Snowflake made Data Science/ML a top priority
- * model training may require *optimized WHs* (expensive)
- * no GPUs → not for NN (but see Container Services w/ NVIDIA)
- * model serving w/ *multi-node standard WHs* (also expensive)
- * is Snowflake ready for Data Science? → still experimental
- * **Snowflake Notebooks**, **Snowflake Cortex** etc coming up

*Check All
Connected Applications*

Tip #95

Embedded Analytics



Checkpoints

- * keep an **architecture diagram** w/ all your integrations
- * understand how they all consume data, use caches...
- * check all in-house developed applications
- * check analytic apps: Tableau, Power BI, Looker...
- * check ETL tools: Fivetran, dbt...

Third-Party Apps
Saving Money
Will Spend Money

Tip #96

Checkpoints

- * they all install in your account → consume YOUR credits!
- * principle: "*spend even more money to save money*"
- * claim: "*we save you more than you spend on us*"
- * they all rely on best practices and tips from this course
- * this proves that Snowflake is not at all a zero-admin DW

Free Marketplace

Native Apps

Will Cost Money

Tip #97

Checkpoints

- * no free Marketplace native app is truly free
- * they all install in your account → use YOUR compute credits
- * plenty of them are actually unreliable → use w/ much care!
- * are better suited for expensive prod account, not test accounts
- * will start an expensive STREAMLIT object
- * may query metadata w/ intensive queries
- * may transparently schedule expensive tasks!

*Keep
App Versions
Updated*

Tip #98

Checkpoints

- * hard to use latest performance and security updates
- * WHs usually have installed older versions → check!
- * check and use the latest deployed framework versions
- * check and use latest deployed packages & libs
- * update your Snowflake client drivers & libraries

*Cache Data
in Third-Party Tools*

Tip #99

Power BI Connection Types



Import Data



Direct Query



Live Connection



Checkpoints

- * mostly for **analytic tools** (Tableau, Power BI, Looker)
- * avoid querying Snowflake data at every visit
- * should combine access to live+cached data
- * use data extracts w/ aggregations for dataviz
- * avoid refreshing dashboards too frequently
- * use query tags to identify intensive operations

*Auto-Abort
Running Queries
from Disconnected Apps*

Tip #100

Checkpoints

- * conn closed on sync queries → all aborted right away
- * conn closed/lost (net outage) on async queries
 - all may continue, for max 2 days!
 - **STATEMENT_TIMEOUT_IN_SECONDS** (account/session/WH level)
 - **ABORT_DETACHED_QUERY** (session level) → abort after 5 mins max



100 Snowflake Cost Optimization Techniques

100 Snowflake Cost Optimization Techniques