

# Interview questions - Caching for Speed Optimization

## ✓ Core Conceptual Questions

---

### ♦ What is caching and why is it important in system design?

#### Answer:

Caching is the technique of storing frequently accessed data in a faster storage layer (e.g., in-memory) to reduce data retrieval time. It helps:

- Minimize latency
  - Reduce backend/database load
  - Improve system scalability and user experience Caching is critical in high-traffic systems where performance and responsiveness are essential.
- 

### ♦ Explain different types of caching and where they are used.

#### Answer:

- **Client-side cache:** Browser cache, service workers — used for static assets, localStorage.
  - **Server-side cache:** In-memory (e.g., Redis, Memcached) — used in API responses or session storage.
  - **CDN cache:** Distributed edge servers cache static content like images, JS, CSS.
  - **Database cache:** Materialized views or query result caching — reduces repeated complex DB queries.
- 

### ♦ What are write-through vs. write-back caching strategies?

#### Answer:

- **Write-through:** Data is written to cache and database simultaneously. Ensures consistency but slower write performance.
  - **Write-back (write-behind):** Data is written to cache first; the DB is updated later asynchronously. Faster writes but at risk of data loss on cache failure.
- 

♦ **What is lazy loading (cache-aside) and when would you use it?**

**Answer: Lazy loading** (cache-aside) loads data into the cache only when needed:

- On a cache miss, fetch from DB → populate cache → return data.
  - Used when not all data is frequently accessed or when cache storage is limited.
  - Gives fine-grained control but requires cache invalidation management.
- 

## **Scenario-Based Questions**

---

♦ **How would you use caching to optimize a product details page?**

**Answer:**

- Cache frequently viewed product data in Redis with a TTL.
  - Use lazy loading so only requested products are cached.
  - Use CDN caching for images.
  - Invalidate or refresh cache on product updates. This reduces DB hits and improves page load speed significantly.
- 

♦ **What eviction strategy would you choose for a memory-limited system?**

**Answer:**

- **LRU (Least Recently Used)** is ideal when recently accessed items are more likely to be used again.
  - **LFU (Least Frequently Used)** if certain items are accessed more often than others.
  - Eviction strategy should match access patterns to avoid cache misses.
- 

♦ **How would you keep cache and database in sync?**

**Answer:**

- Use **write-through** for strong consistency.
  - Use **cache invalidation** on data updates (manually or via message queues).
  - Set appropriate TTLs to auto-refresh stale data.
  - Optionally, use change data capture (CDC) mechanisms or event-driven updates.
- 

♦ **What are the potential downsides or risks of aggressive caching?**

**Answer:**

- **Stale data:** Cached values may be outdated if not invalidated properly.
  - **Cache stampede:** Multiple requests for uncached data can hit the backend simultaneously.
  - **Overuse of memory:** Poor eviction strategy can lead to inefficient memory usage.
  - **Complexity:** Cache invalidation and consistency handling can increase system complexity.
- 

## **Practical Implementation**

---

♦ **How would you implement Redis caching in a web application?**

**Answer:**

- Use Redis as a key-value store (e.g., `product:123` → `productData`).
  - On cache miss, fetch from DB, store in Redis with TTL.
  - On cache hit, return directly from Redis.
  - Use libraries like StackExchange.Redis (C#), `redis-py` (Python), or `ioredis` (Node.js).
- 

### ♦ How can you prevent cache stampedes or thundering herd problems?

Answer:

- Use **lock or mutex** during cache miss to ensure only one backend call populates the cache.
  - Use **stale-while-revalidate** strategies to serve old data while refreshing in the background.
  - Implement **randomized TTLs** to spread out cache expiry.
- 

### ♦ What tools can be used for distributed caching?

Answer:

- **Redis Cluster**: High availability and partitioning.
- **Memcached**: Lightweight in-memory store.
- **Hazelcast, Apache Ignite**: In-JVM distributed cache for enterprise use.
- Use consistent hashing and replication to manage cache distribution across nodes.