Interview questions - Messaging & Queues for Decoupling

✓ 1. Why would you use asynchronous messaging in a system?

Answer: Asynchronous messaging decouples the sender and receiver, allowing them to operate independently. This improves:

- **Scalability**: Producers don't wait for consumers.
- **Resilience**: If the consumer is slow or down, messages are queued.
- **Performance**: Reduces response time for users by offloading background tasks (e.g., sending emails, processing orders).
- Loose coupling: Systems can evolve independently.

Example: Sending a confirmation email after a purchase — no need to delay user response for email delivery.

2. What are the differences between RabbitMQ and Kafka? When would you use one over the other?

Answer:

Feature	RabbitMQ	Kafka
Message Model	Queue-based (push)	Log-based (pull)
Ordering	FIFO per queue	Per-partition ordering
Retention	Deletes message once consumed	Retains for configurable time
Use Case Fit	Real-time, transactional workloads	High-throughput data pipelines, event sourcing
Delivery Guarantee	At-most-once / At-least-once	At-least-once / Exactly-once (with config)

When to use:

- Use RabbitMQ for complex routing, low-latency tasks, and per-message acknowledgment (e.g., order processing, real-time updates).
- Use **Kafka** for **event streaming**, **log aggregation**, and **analytics** (e.g., tracking user behavior at scale).

☑ 3. Explain the different message delivery guarantees: At-least-once vs. Exactly-once vs. At-most-once – what are they and where do you apply them?

Answer:

- At-most-once: Message is sent once and not retried on failure.
 - ➤ Fast but can lose data. Used in non-critical logging.
- At-least-once: Message is retried until acknowledged.
 - ➤ Default in most systems. May result in duplicates (use with idempotency).
- **Exactly-once**: Message is delivered once with no duplication.
 - > Complex and expensive. Used in critical financial transactions.

Use Cases:

- **Exactly-once**: Bank transfers, billing systems.
- At-least-once: Order placement, email dispatching.
- At-most-once: Non-critical telemetry/logs.

✓ 4. How do queues help improve system scalability and fault tolerance?

Answer:

- **Scalability**: Queues act as a buffer. Producers send messages quickly, and consumers can scale independently to process them.
- **Fault Tolerance**: If a consumer crashes, messages remain in the queue until it recovers or another consumer takes over.

• Load Management: Prevents consumer overload; helps smooth traffic spikes.

Example: Spiking traffic on Black Friday — queue absorbs the spike and allows backend systems to catch up.

5. What problems might arise in a queue-based system under heavy load? How would you mitigate them?

Answer: Problems:

- Queue backlog: Slower consumers can't keep up.
- Message delay: Increased latency.
- Message duplication (if retries occur).
- Resource exhaustion: CPU/memory strain from backlogs.

Mitigations:

- Scale consumers horizontally (auto-scaling).
- Optimize message processing logic.
- Use dead-letter queues to isolate poison messages.
- Implement rate limiting and backpressure mechanisms.

6. How would you design an order processing system using a message queue?

Answer:

- 1. **User places order** → Producer service sends message to queue.
- 2. Order Service (consumer):
 - o Reads from queue.
 - Validates stock.

- o Reserves inventory.
- o Triggers downstream services: payment, confirmation email.
- 3. Each step may publish its own events (event-driven chaining).

Queue helps by:

- Decoupling payment/inventory from frontend latency.
- Handling retries on failure.
- Scaling services independently.

7. How would you ensure idempotency in the consumers?

Answer:

- Use unique message IDs (e.g., Order ID, Message ID).
- Maintain a processed-message store (e.g., Redis or DB) to track already-handled messages.
- Make operations idempotent:
 - o E.g., "Mark order as processed" only if not already marked.
- Leverage deduplication features if provided by the message broker.

Why? Prevents duplicate processing in **at-least-once delivery** scenarios.